

# ISA 444: Business Forecasting

## 02: Introduction to

Fadel M. Megahed, PhD

Endres Associate Professor  
Farmer School of Business  
Miami University

 @FadelMegahed

 fmegahed

 fmegahed@miamioh.edu

 Automated Scheduler for Office Hours

Spring 2023

# Quick Refresher from Last Class

- Describe **course motivation** and **structure**.
- Explain the differences between **cross sectional**, **time series** and **panel** datasets.
- Describe the **components of time series** datasets.
- Explain the **forecasting steps**.

# Learning Objectives for Today's Class

- Describe the syntax, data types, and data structures in `R`.
- Access the help for `R` functions (each help file has the following components: Description, usage, arguments, value, and examples).
- Utilize the project workflow in `R` and create `R` script.
- Access, subset, and create `ts()` objects in `R`.

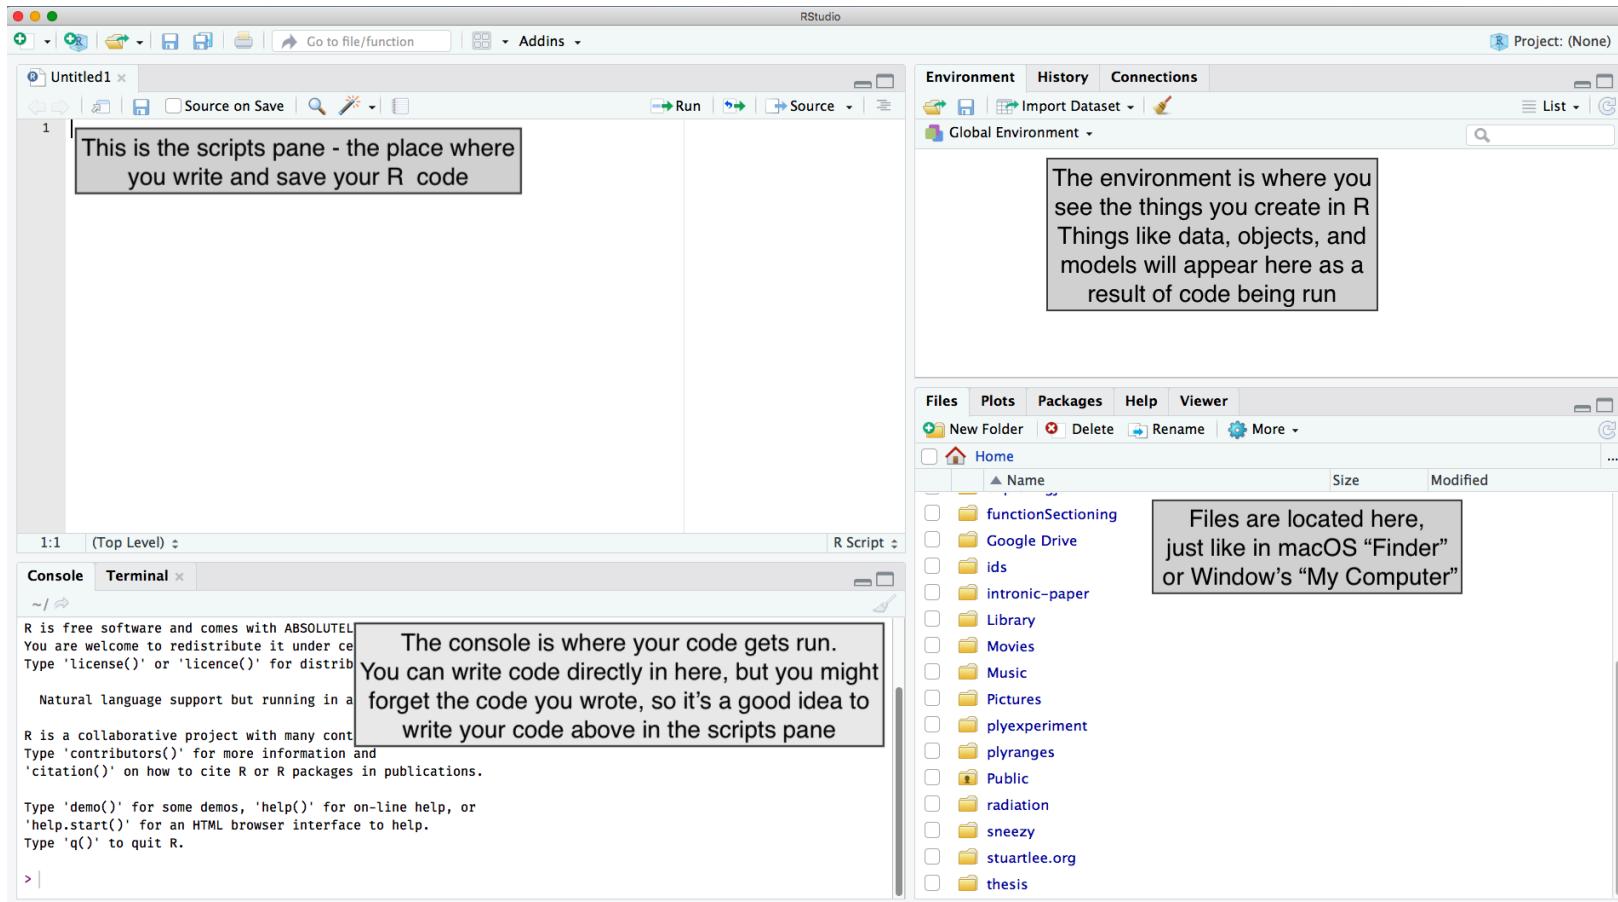
# Learning (Any Programming Language)



-  **Get hands dirty!!**
-  Documentation! Documentation! Documentation!
-  (Not surprisingly) Learn to Google: what that error message means (I **G**a lot 😅)

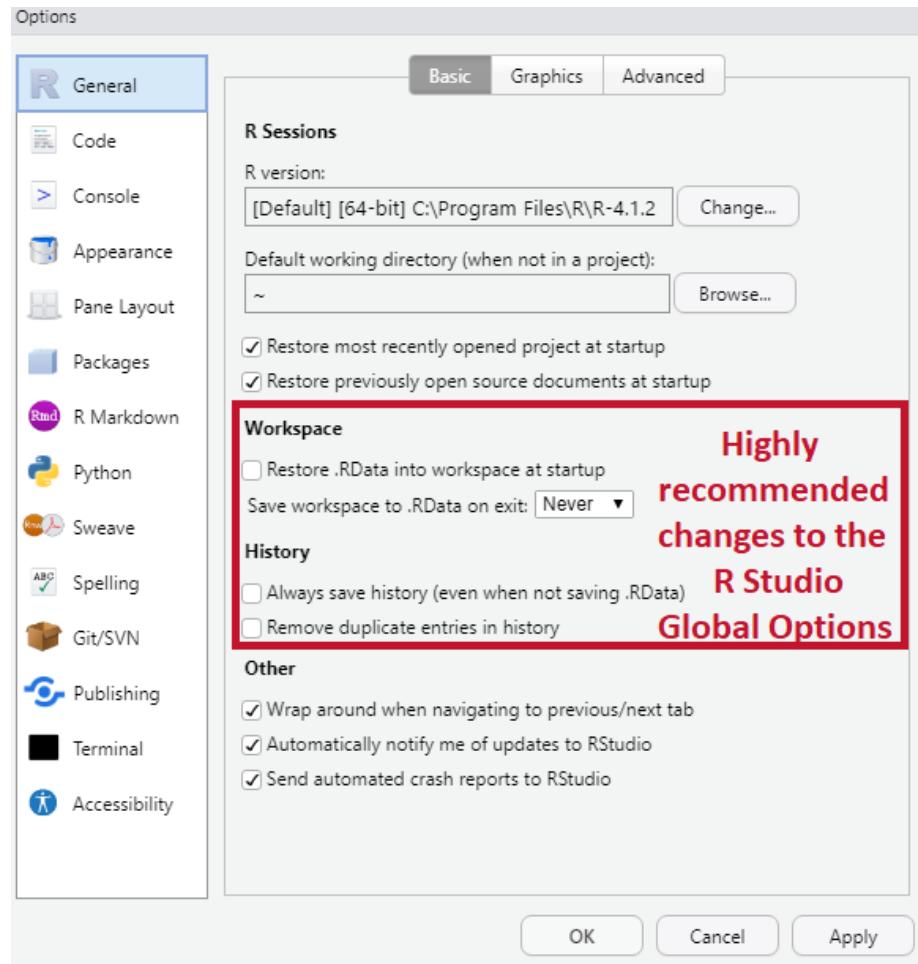
# The RStudio Interface, Setup and a Project-Oriented Workflow for your Analysis

# RStudio Interface



# Setting up RStudio (do this once)

Go to **Tools > Global Options**:



Uncheck **Workspace** and **History**, which helps to keep  working environment fresh and clean every time you switch between projects.

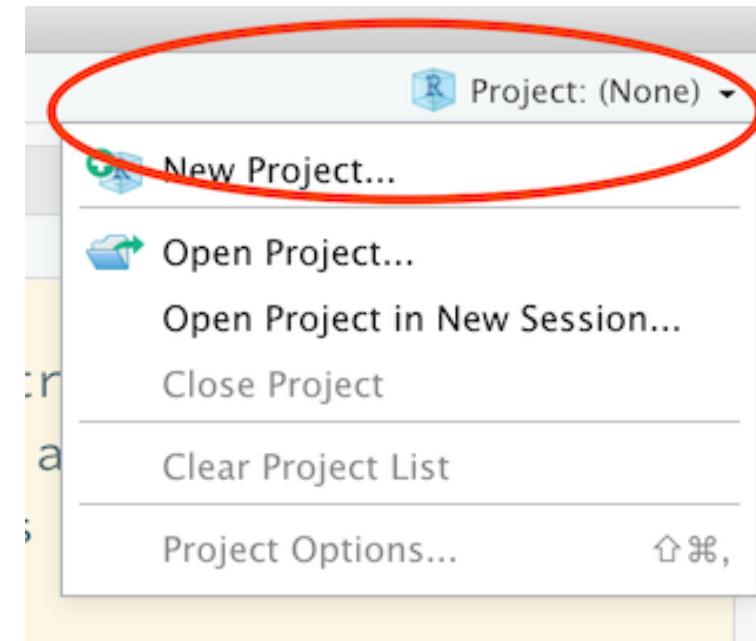
# What is a Project?

- Each university course is a project, and get your work organised.
- A self-contained project is a folder that contains all relevant files, for example my `ISA 444/`  
 includes:
  - `isa444.Rproj`
  - `lectures/`
    - `01_introduction/`
      - `01-intro.Rmd`, etc.
    - `02_introduction_to_r/`
      - `02_intro_r.Rmd`, etc.
  - All working files are **relative** to the **project root** (i.e. `isa444/`).
  - The project should just work on a different computer.

02:00

# Lets Create a .Rproj for Our Course

- Click the **Project** icon on the top right corner
- **New Directory/Existing Directory > New Project > Create Project**
- Open the project





# 101: Syntax, Data Types, Data Structures and Functions

# Coding Style

Good coding style is like correct punctuation: you can manage without it, but it sure makes things easier to read.

— [The tidyverse style guide](#)

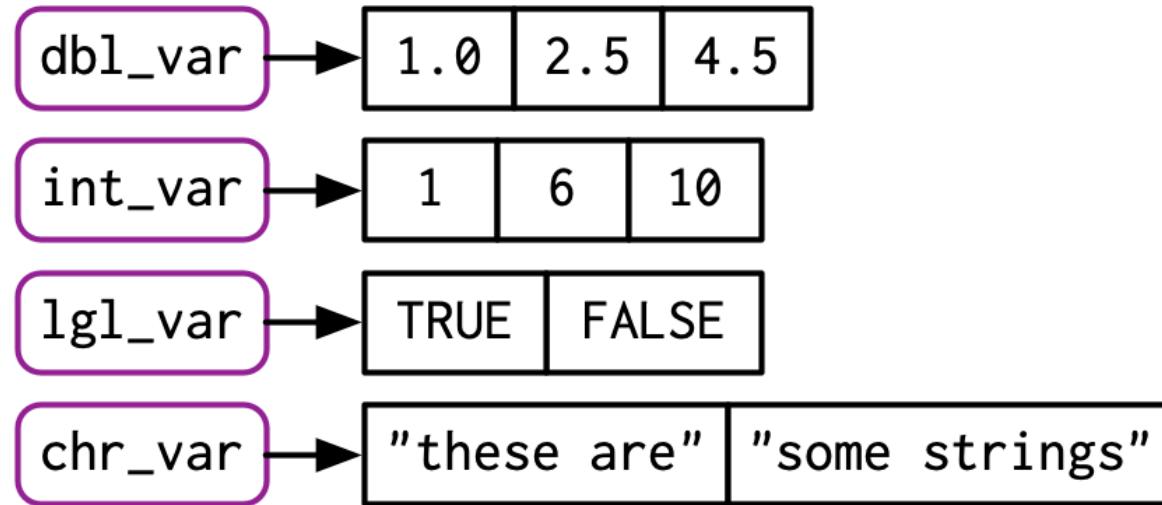
## R style guide

✓ snake\_case

✗ camelCase (Javascript)

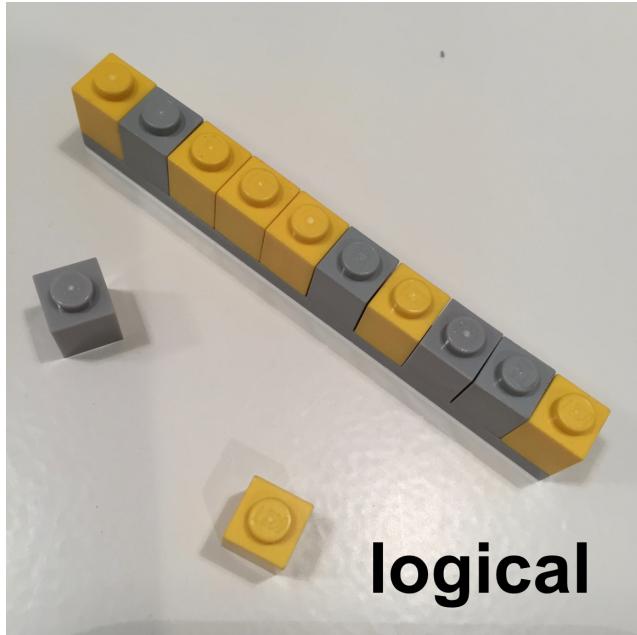
✗ PascalCase (Python)

# Data Types: A Visual Introduction

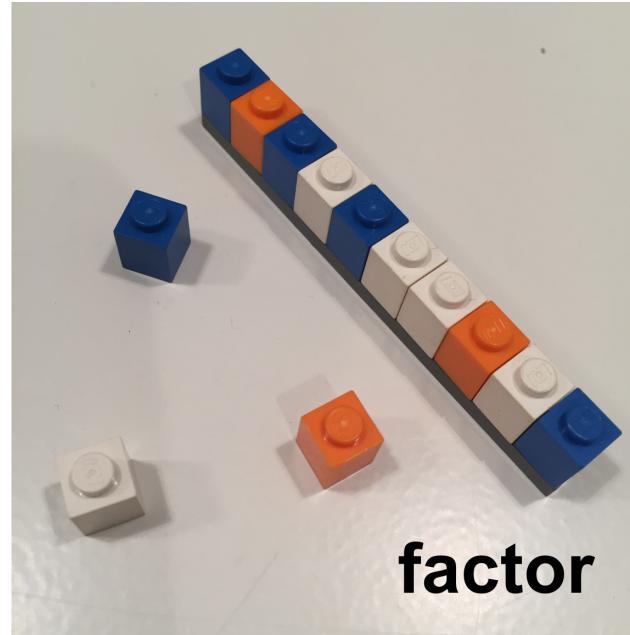


- To check the **type of** an object in , you can use the function **typeof**.

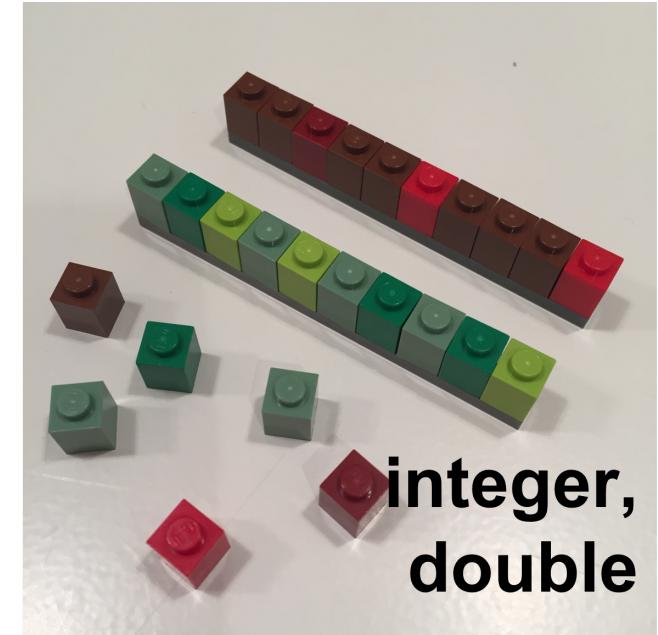
# Data Types: A Visual Introduction



**logical**



**factor**



**integer,  
double**

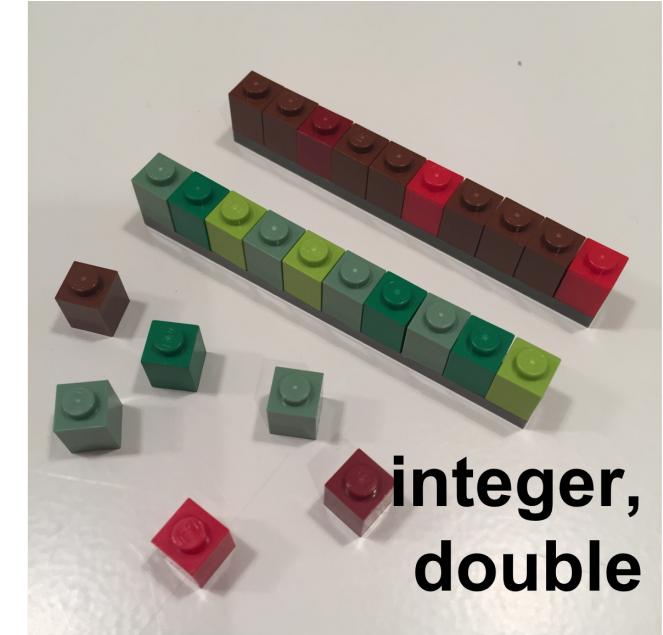
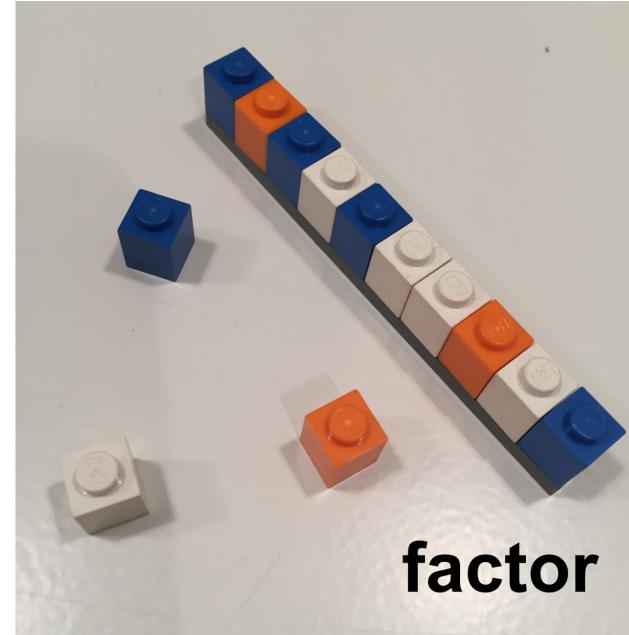
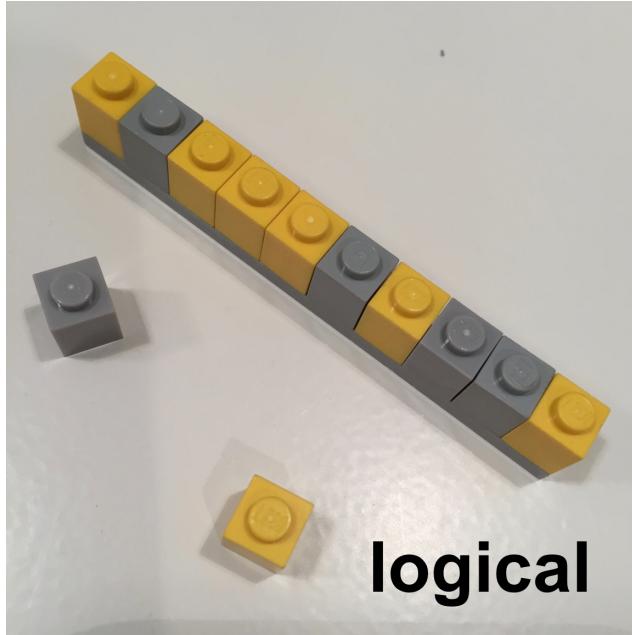
A visual representation of different types of atomic vectors

# Data Types: Formal Definitions

Each of the four primary types has a special syntax to create an individual value:

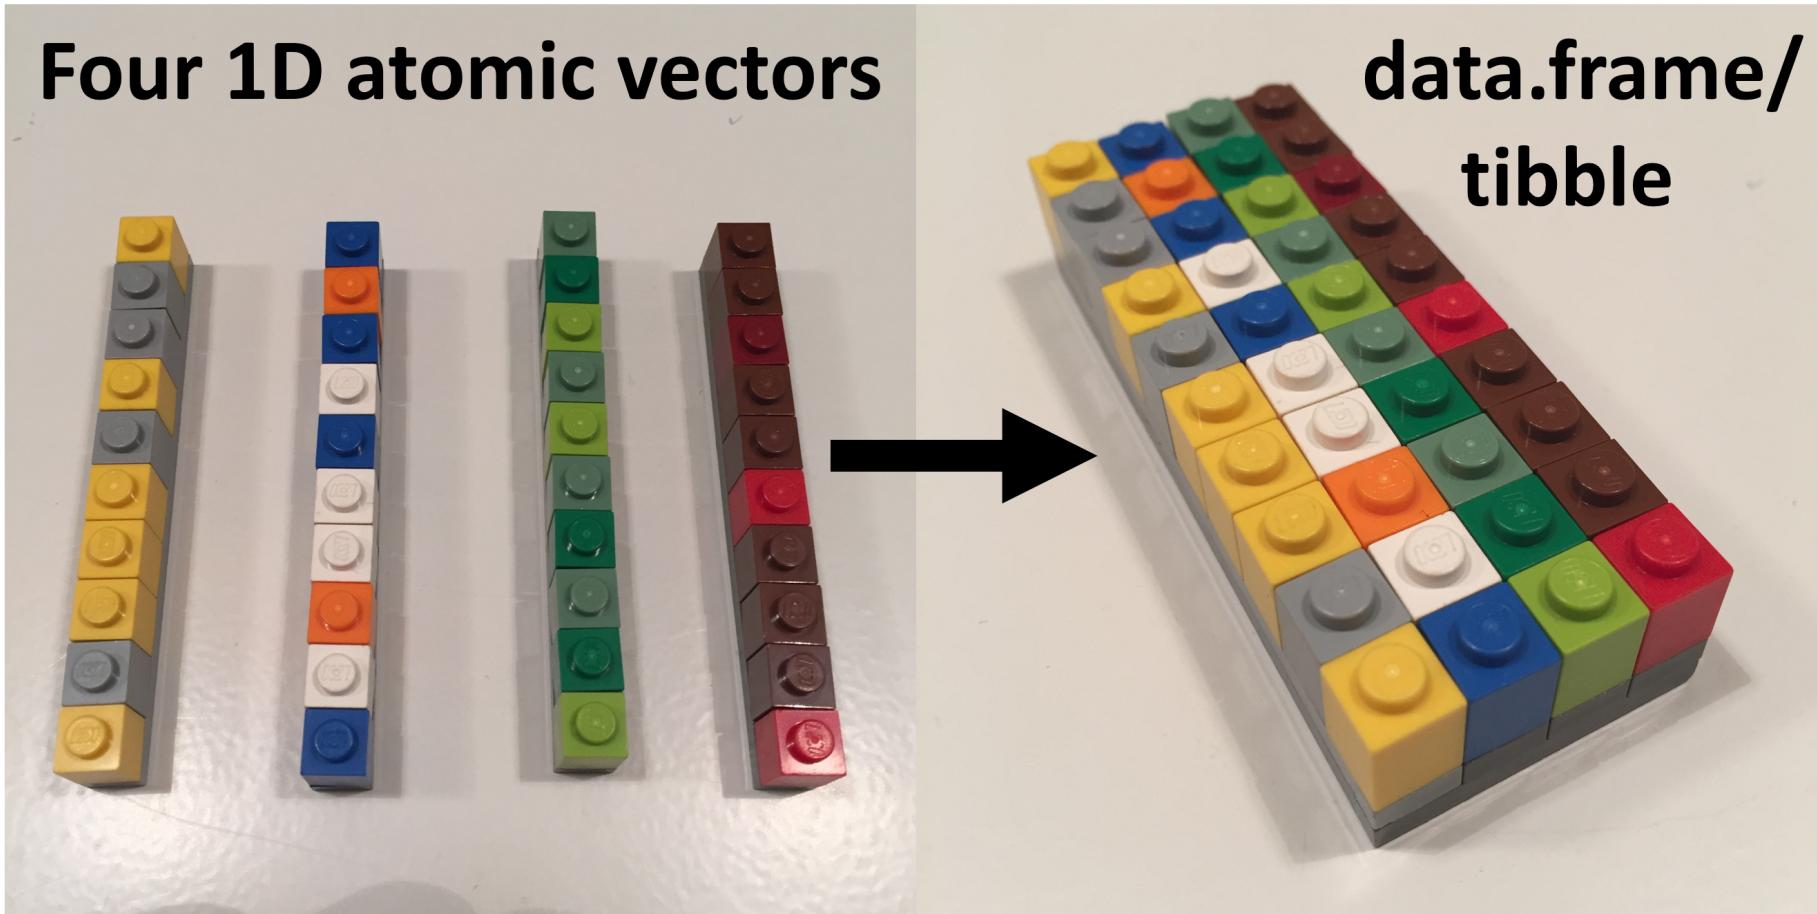
- Logicals can be written in full (`TRUE` or `FALSE`), or abbreviated (`T` or `F`).
- Doubles can be specified in decimal (`0.1234`), scientific (`1.23e4`), or hexadecimal (`0xcafe`) form.
  - There are three special values unique to doubles: `Inf`, `-Inf`, and `NaN` (not a number).
  - These are special values defined by the floating point standard.
- Integers are written similarly to doubles but must be followed by `L` (`1234L`, `1e4L`, or `0xcafeL`), and can not contain fractional values.
- Strings are surrounded by `"` (e.g., `"hi"`) or `'` (e.g., `'bye'`). Special characters are escaped with `\` see `?Quotes` for full details.

# Data Structures: Atomic Vector (1D)



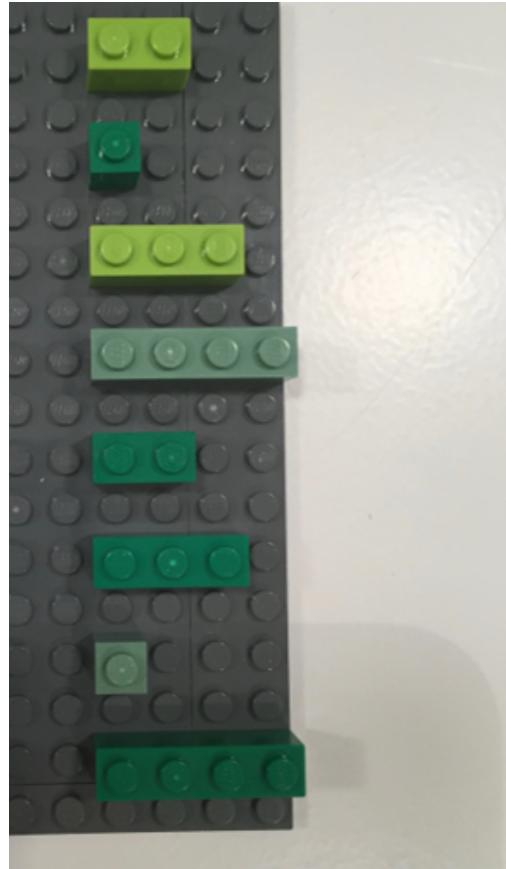
Keeping the visual representation of different types of atomic vectors in your head!!

# Data Structures: 1D → 2D



# Data Structures: Lists

An object contains elements of **different data types**.



# Data Structures: Lists



```
lst <- list( # list constructor/creator
  1:3, # atomic double/numeric vector of length = 3
  "a", # atomic character vector of length = 1 (aka scalar)
  c(TRUE, FALSE, TRUE), # atomic logical vector of length = 3
  c(2.3, 5.9) # atomic double/numeric vector of length =3
)
lst # printing the list
```

```
## [1] "1:3"                      "a"                         "c(TRUE, FALSE, TRUE)"
## [4] "c(2.3, 5.9)"
```

# Data Structures: Lists

Subset by []

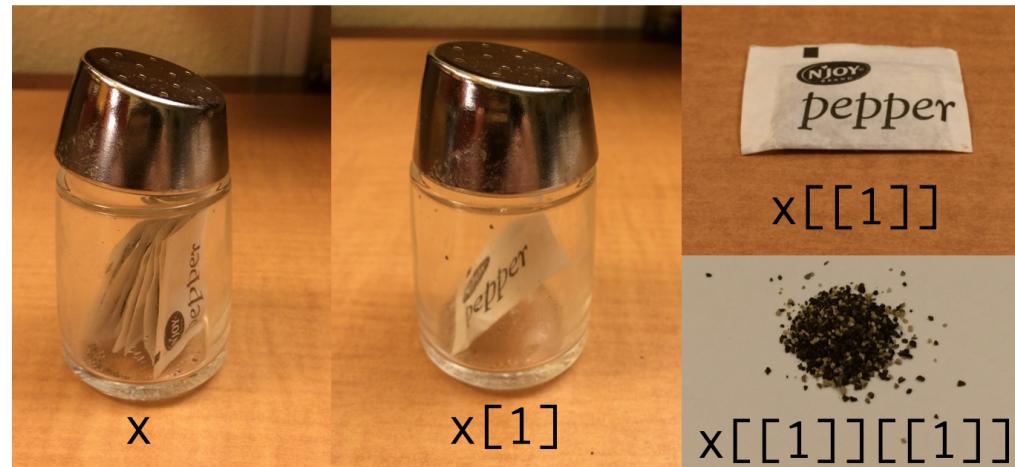
```
lst[1]
```

```
## [[1]]  
## [1] 1 2 3
```

Subset by [[]]

```
lst[[1]]
```

```
## [1] 1 2 3
```



# Data Structures: Matrices

A matrix is a **2D data structure** made of **one/homogeneous data type**.

```
x_mat = matrix( sample(1:10, size = 4), n  
str(x_mat) # its structure?
```

```
## int [1:2, 1:2] 7 5 1 6
```

```
x_mat # printing it nicely  
print('-----')  
x_mat[1, 2] # subsetting
```

```
##      [,1] [,2]  
## [1,]    7    1  
## [2,]    5    6  
## [1] "-----"  
## [1] 1
```

# Data Structures: Matrices

A matrix is a **2D data structure** made of **one/homogeneous data type**.

```
x_mat = matrix( sample(1:10, size = 4), n  
str(x_mat) # its structure?
```

```
## int [1:2, 1:2] 7 5 1 6
```

```
x_mat # printing it nicely  
print('-----')
```

```
x_mat[1, 2] # subsetting
```

```
##      [,1] [,2]  
## [1,]    7    1  
## [2,]    5    6  
## [1] "-----"  
## [1] 1
```

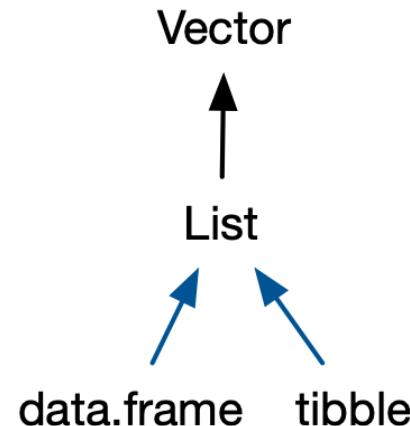
```
x_char = matrix(  
  sample(letters, size = 12), nrow = 3, n  
x_char
```

```
##      [,1] [,2] [,3] [,4]  
## [1,] "d"  "e"  "a"  "u"  
## [2,] "v"  "s"  "k"  "z"  
## [3,] "b"  "n"  "l"  "t"
```

```
x_char[1:2, 2:3] # subsetting
```

```
##      [,1] [,2]  
## [1,] "e"  "a"  
## [2,] "s"  "k"
```

# Data Structures: Data Frames



If you do data analysis in R, you're going to be using data frames. A data frame is a named list of vectors with attributes for `(column) names`, `row.names`, and its class, “`data.frame`”. -- Hadley Wickham

# Data Structures: Data Frames

```
df1 <- data.frame(x = 1:3, y = letters[1:3])
typeof(df1) # showing that its a special case of a list
```

```
## [1] "list"
```

```
attributes(df1) # but also is of class data.frame
```

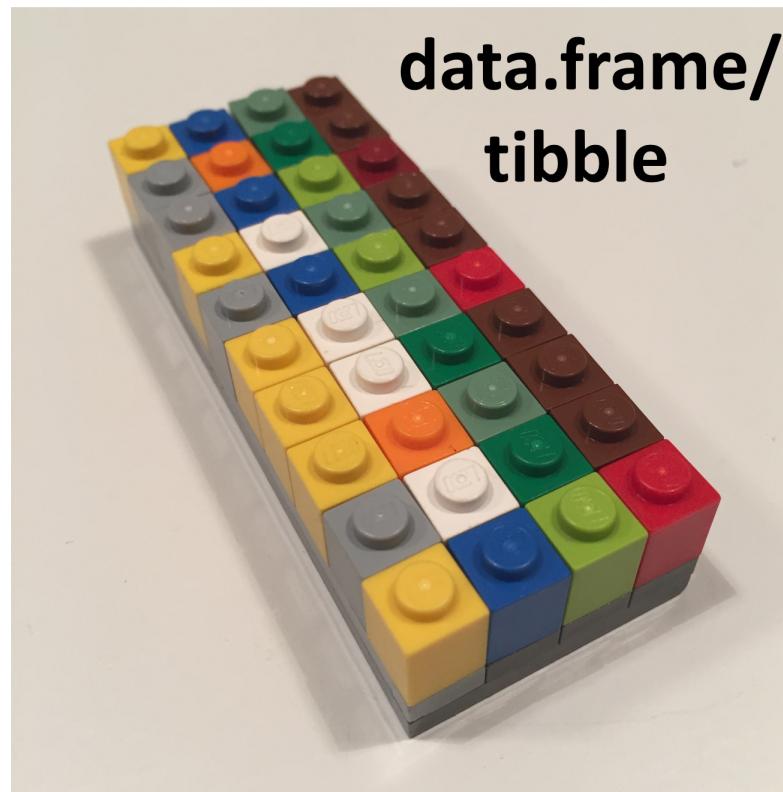
```
## $names
## [1] "x" "y"
##
## $class
## [1] "data.frame"
##
## $row.names
## [1] 1 2 3
```

In contrast to a regular list, a data frame has **an additional constraint: the length of each of its vectors must be the same**. This gives data frames their **rectangular structure**.

Source: Content is from Hadley Wickham's Advanced R: Chapter 3 on Vectors

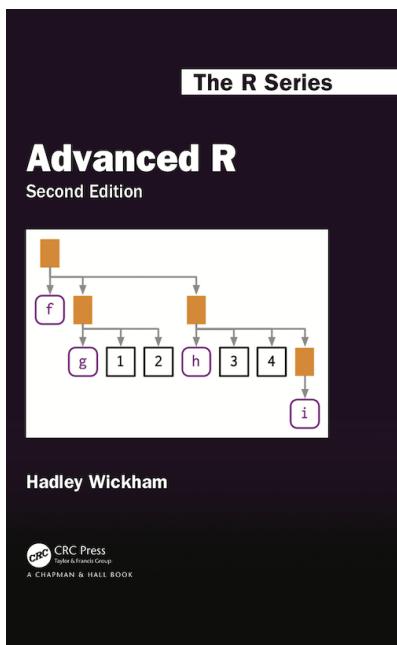
# Data Structures: Data Frames

As noted in the creation of `df1`, columns in a data frame can be of different types. Hence, it is more widely used in data analysis than matrices.



# Data Structures: So What is a Tibble Anyway?

Tibble is a **modern reimagining of the data frame**. Tibbles are designed to be (as much as possible) **drop-in replacements for data frames** that fix those frustrations. A concise, and fun, way to summarise the main differences is that tibbles are **lazy and surly: they do less and complain more**. -- Hadley Wickham



To learn more about the basics of tibble, please consult the reference below:

- Data frames and tibbles (Click and read from 3.6 up to and including 3.6.5)

# Functions

A function call consists of the **function name** followed by one or more **argument** within parentheses.

```
temp_high_forecast = c(34, 36, 41, 44, 27, 32, 35)
mean(x = temp_high_forecast)
```

```
## [1] 35.57143
```

- function name: `mean()`, a built-in R function to compute mean of a vector
- argument: the first argument (LHS `x`) to specify the data (RHS `temp_high_forecast`)

# Function Help Page

Check the function's help page with `?mean`

## Class Activity

*Please take 2 minutes to investigate the help page for `mean` in R Studio.*

```
mean(x = temp_high_forecast, trim = 0, na.rm = FALSE, ...)
```

- Read **Usage** section
  - What arguments have default values?
- Read **Arguments** section
  - What does `trim` do?
- Run **Example** code

# Function Arguments

## Match by positions

```
mean(temp_high_forecast, 0.1, TRUE)
```

```
## [1] 35.57143
```

## Match by names

```
mean(x = temp_high_forecast, trim = 0.1,
```

```
## [1] 35.57143
```

# Use Functions from Packages

```
library(dplyr)  
dplyr::cummean(temp_high_forecast)
```

```
## [1] 34.00000 35.00000 37.00000 38.7500
```

```
dplyr::first(temp_high_forecast)
```

```
## [1] 34
```

```
dplyr::last(temp_high_forecast)
```

```
## [1] 35
```

```
install.packages("light")
```



```
library("light")
```



Images sourced from <https://www.wikihow.com/Change-a-Light-Bulb>

# R for time series analysis

# The *ts* Object

R facilitates the work with time series data by storing it in an aptly named ts object.

# Create a *ts* Object by asking *ChatGPT*

The screenshot shows the ChatGPT interface. At the top, there's a header bar with a user profile picture, the title "Creating a ts object in R", and various interaction icons like a message count (0), a link icon, and a share icon. Below the header, there's a timestamp ("1 min") and a view count ("0 views"). On the left side, there's a sidebar with options like "Control Charts In SPC", "Clear conversations", "Dark mode", "OpenAI Discord", and "Updates & FAQ". The main area is titled "ChatGPT" and contains three sections: "Examples", "Capabilities", and "Limitations". A large play button is centered over the "Capabilities" section. The "Examples" section shows two examples of user input and AI output. The "Capabilities" section lists the AI's ability to remember conversation history, provide follow-up corrections, and decline inappropriate requests. The "Limitations" section notes that the AI may generate incorrect information, produce harmful instructions, and have limited knowledge of world events after 2021. At the bottom, there's a footer with the text "ChatGPT Jan 9 Version. Free Research Preview. Our goal is to make AI systems more natural and safe to interact with. Your feedback will help us improve." and a copyright notice "© 2023 OpenAI Inc. All rights reserved."

# Demo 1: Nile & AirPassengers Data

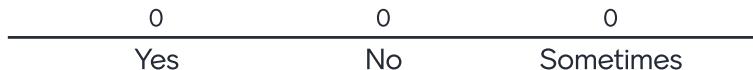
In class, we will create our first  script where we will examine one of these two built-in datasets. In our exploration, we will:

- Examine the `typeof()` the dataset.
- Examine the `class()` of the dataset.
- Examine the `length()` of the dataset.
- `print()` the data set and examine its `frequency()`.
- Subset the data using `window()` and non ts-based sub-setting techniques.
  - Useful for the concept of the entire time series vs a snippet that we discussed in [Class 01](#).

# A Poll: File Paths in R

Go to [www.menti.com/alyzfp5phujn](http://www.menti.com/alyzfp5phujn)

Do you often struggle with identifying the file path  
when reading data in R?



```
```{r}
#Chunk 1
library(ggplot2)
library(xlsx)
library(dplyr)
library(dplyr)

SRF <- read_excel("Desktop/StudentRevertantFrequencies.xlsx")
```

Error: path does not exist:
'Desktop/StudentRevertantFrequencies.xlsx'
 3. stop("`path` does not exist: ", sQuote(path), call. = FALSE)
 2. check_file(path)
 1. read_excel("Desktop/StudentRevertantFrequencies.xlsx")
```



# Demo 2: Loading Time Series Data into R

- Reading CSV Files through either the `read.csv()` or `readr::read_csv()` functions.
- Reading FRED and Yahoo Finance Data using the `tidyquant` package.
- Preprocessing and converting the data into a time series object.

## Bonus Tricks (if time allows)

- `file.choose()` -- not typically introduced by Miami faculty!!
- The datapasta package 

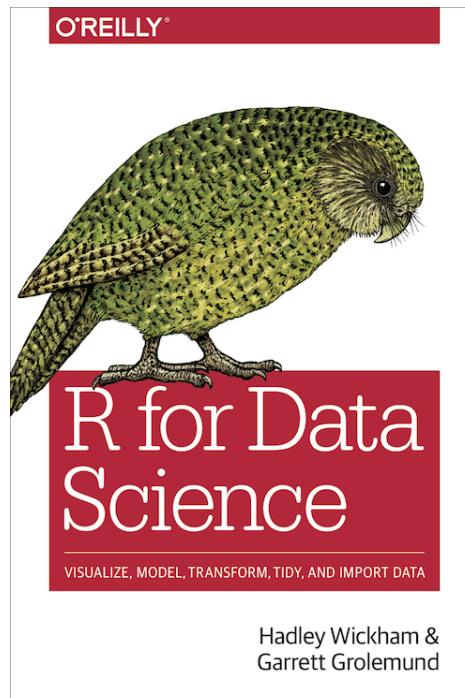
# Summary of Main Points

By now, you should be able to do the following:

- Describe the syntax, data types, and data structures in `R`.
- Access the help for `R` functions (each help file has the following components: Description, usage, arguments, value, and examples).
- Utilize the project workflow in `R` and create `R` script.
- Access, subset, and create `ts()` objects in `R`.

# Things to Do to Prepare for Our Next Class

- Go over your notes, read the **references below**, and **complete** the self-paced R tutorial.
- Complete [Assignment 02](#) on Canvas.



- Data Visualization
- Graphics for Communication
- Dates and Times