

# ISA 444: Business Forecasting

## 06: The Naive Forecast, Measures of Forecast Accuracy and the Prediction Interval

Fadel M. Megahed, PhD

Endres Associate Professor  
Farmer School of Business  
Miami University

 @FadelMegahed

 fmegahed

 fmegahed@miamioh.edu

 Automated Scheduler for Office Hours

Spring 2023

# Quick Refresher from Last Class

- ✓ Use numerical summaries to describe a time series.
- ✓ Explain what do we mean by correlation.
- ✓ Apply transformations to a time series.

# Recap: Viz + Numerical Summary = Big Picture


set ▾	x.mean ▾	x.sd ▾	y.mean ▾	y.sd ▾	corr ▾
I	9	3.32	7.5	2.03	0.82
II	9	3.32	7.5	2.03	0.82
III	9	3.32	7.5	2.03	0.82
IV	9	3.32	7.5	2.03	0.82

Showing 1 to 4 of 4 entries

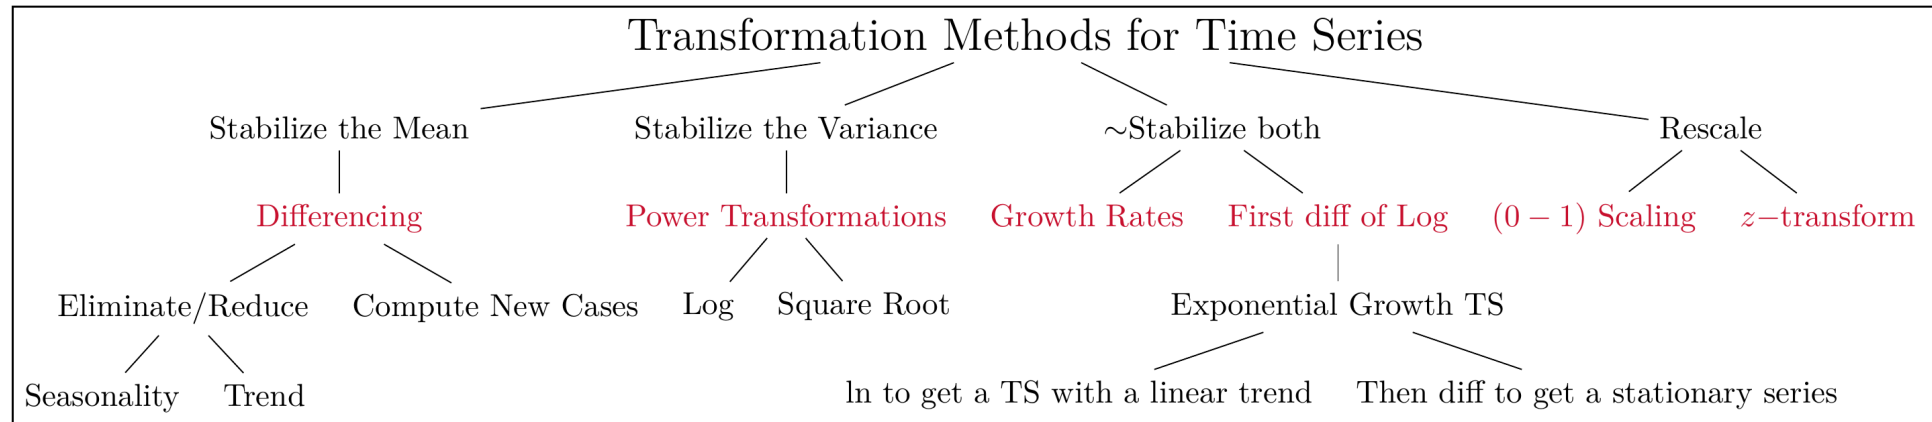
Previous

1

Next



# Recap: Guidelines for Transforming TS Data



A classification of common transformation approaches for time series data

# A Quick Recap of Assignment 04

## Question 01

```
tsla = tidyquant::tq_get(  
  x = 'tsla', from = '2020-01-01', to = Sys.Date(),  
  periodicity = 'monthly'  
) |>  
  dplyr::mutate(  
    year = lubridate::year(date) |> as.factor(),  
    month = lubridate::month(date, label = T)  
  )  
tsla
```

```
## # A tibble: 38 × 10  
##   symbol date      open high low  
##   <chr> <date>    <dbl> <dbl> <dbl>  
## 1 tsla  2020-01-01  28.3  43.5  28.1  
## 2 tsla  2020-02-01  44.9  64.6  40.8  
## 3 tsla  2020-03-01  47.4  53.8  23.4  
## 4 tsla  2020-04-01  33.6  58.0  29.8  
## 5 tsla  2020-05-01  50.3  56.2  45.5  
## 6 tsla  2020-06-01  57.2  72.5  56.9  
## 7 tsla  2020-07-01  72.2 120.   72.6  
## 8 tsla  2020-08-01  96.6 167.   91  
## 9 tsla  2020-09-01 167. 167. 110.  
## 10 tsla 2020-10-01 147. 155. 126.  
## # ... with 28 more rows
```

# A Quick Recap of Assignment 04

## Question 01 ❌

```
tsla |>
  ggplot2::ggplot(
    ggplot2::aes(x = month, y = adjusted, color = year)
  ) + # aesthetics
  ggplot2::geom_point() + # adding the points
  ggplot2::geom_line() + # adding a line for each year
  ggplot2::scale_y_continuous(breaks = scales::pretty_breaks(n = 6))
  ggplot2::scale_color_brewer(palette = 'Dark2') + # color blind friendly
  ggplot2::labs(
    x = 'Month',
    y = 'Adjusted Closing Price',
    title = 'The adjusted closing price of TSLA'
  ) +
  ggplot2::theme_bw() + # black and white theme
  ggplot2::theme(legend.position = 'bottom') # put legend below chart
```




# A Quick Recap of Assignment 04

## Question 01




```
tsla |>
  ggplot2::ggplot(
    ggplot2::aes(x = month, y = adjusted, color = year, group = year)
  ) + # aesthetics
  ggplot2::geom_point() + # adding the points
  ggplot2::geom_line() + # adding a line for each year
  ggplot2::scale_y_continuous(breaks = scales::pretty_breaks(n = 6))
  ggplot2::scale_color_brewer(palette = 'Dark2') + # color blind friendly
  ggplot2::labs(
    x = 'Month',
    y = 'Adjusted Closing Price',
    title = 'The adjusted closing price of TSLA'
  ) +
  ggplot2::theme_bw() + # black and white theme
  ggplot2::theme(legend.position = 'bottom') # put legend below chart
```



# A Quick Recap of Assignment 04

## Question 02

```
tsla |>
  ggplot2::ggplot(
    ggplot2::aes(x = date, y = adjusted, color = year, group = year)
  ) + # aesthetics
  ggplot2::geom_point() + # adding the points
  ggplot2::geom_line() + # adding a line for each year
  ggplot2::scale_y_continuous(breaks = scales::pretty_breaks(n = 6))
  ggplot2::scale_color_brewer(palette = 'Dark2') + # color blind friendly
  ggplot2::labs(
    x = 'Month',
    y = 'Adjusted Closing Price',
    title = 'The adjusted closing price of TSLA'
  ) +
  ggplot2::theme_bw() + # black and white theme
  ggplot2::theme(legend.position = 'bottom') # put legend below chart
```





# A Quick Recap of Assignment 04

## Question 03

This will obviously depend on your stock. In my case, there is no evidence for a consistent seasonal pattern.

# A Quick Recap of Assignment 04

## Question 04

In question 4, the mean makes sense to compute. The standard deviation does not since you only have one observation per month.

```
tsla_summary =  
  tsla |>  
  dplyr::group_by(symbol, year, month) |>  
  dplyr::summarise(  
    adj_avg = mean(adjusted),  
    adj_sd = sd(adjusted)  
  )  
  
print(tsla_summary, n = 10)
```

```
## # A tibble: 38 × 5  
## # Groups:   symbol, year [4]  
##   symbol year month adj_avg adj_sd  
##   <chr>  <fct> <ord>   <dbl> <dbl>  
## 1  tsla  2020  Jan     43.4    NA  
## 2  tsla  2020  Feb     44.5    NA  
## 3  tsla  2020  Mar     34.9    NA  
## 4  tsla  2020  Apr     52.1    NA  
## 5  tsla  2020  May     55.7    NA  
## 6  tsla  2020  Jun     72.0    NA  
## 7  tsla  2020  Jul     95.4    NA  
## 8  tsla  2020  Aug    166.    NA  
## 9  tsla  2020  Sep    143.    NA  
## 10 tsla  2020  Oct    129.    NA  
## # ... with 28 more rows
```

# A Quick Recap of Assignment 04

## Question 05

```
nasdaq =
  tidyquant::tq_get(x = '^IXIC', from = '2020-01-01', to = Sys.Date())
  periodicity = 'monthly') |>
  dplyr::mutate(
    year = lubridate::year(date) |> as.factor(),
    month = lubridate::month(date, label = T)
  )

# putting them next to each other (note the adjusted column name char
both_stocks = dplyr::left_join(
  x = tsla |> dplyr::select(date, year, month, adjusted),
  y = nasdaq |> dplyr::select(date, adjusted),
  by = 'date'
)

both_stocks
```

```
## # A tibble: 38 × 5
##   date          year month adjusted.x a
##   <date>        <fct> <ord>      <dbl>
## 1 2020-01-01  2020  Jan        43.4
## 2 2020-02-01  2020  Feb        44.5
## 3 2020-03-01  2020  Mar        34.9
## 4 2020-04-01  2020  Apr        52.1
## 5 2020-05-01  2020  May        55.7
## 6 2020-06-01  2020  Jun        72.0
## 7 2020-07-01  2020  Jul        95.4
## 8 2020-08-01  2020  Aug       166.
## 9 2020-09-01  2020  Sep       143.
## 10 2020-10-01 2020  Oct       129.
## # ... with 28 more rows
```

# A Quick Recap of Assignment 04

## Question 05

```
nasdaq =
  tidyquant::tq_get(x = '^IXIC', from = '2020-01-01', to = Sys.Date())
  periodicity = 'monthly') |>
  dplyr::mutate(
    year = lubridate::year(date) |> as.factor(),
    month = lubridate::month(date, label = T)
  )

# putting them next to each other (note the adjusted column name char
both_stocks = dplyr::left_join(
  x = tsla |> dplyr::select(date, year, month, adjusted),
  y = nasdaq |> dplyr::select(date, adjusted),
  by = 'date'
)

both_stocks |>
  dplyr::group_by(year) |>
  dplyr::summarise(corr = cor(adjusted.x, adjusted.y))
```

```
## # A tibble: 4 × 2
##   year    corr
##   <fct> <dbl>
## 1 2020  0.959
## 2 2021  0.721
## 3 2022  0.792
## 4 2023   1
```

# Learning Objectives for Today's Class

- Compute the nonseasonal naive forecast.
- Apply and interpret measures of forecast accuracy.
- Interpret prediction intervals for a simple forecast.

# The Naive Forecast

# Recap: What is Forecasting?

[ALL](#) [NEWS](#) [IMAGES](#) [VIDEOS](#) [MAPS](#) [SHOPPING](#) [MORE](#)


About 16,900,000 results [Any time](#) ▾

Dictionary

Data from [Oxford Languages](#)

[Look it up](#)

## fore·cast

[ˈfɔːr,kɑːst] 

**VERB**

1. predict or estimate (a future event or trend):  
*"rain is forecast for eastern Ohio" · "coal consumption is forecast to increase"*

SIMILAR: [predict](#) [prophecy](#) [prognosticate](#) [augur](#) [divine](#) [foretell](#) ▾

---

**NOUN**


1. a prediction or estimate of future events, especially coming weather or a financial trend.

SIMILAR: [prediction](#) [prophecy](#) [forewarning](#) [prognostication](#) [augury](#) ▾

Translate forecast to

[Choose language](#) ▾

[More definitions and word origin](#)



# A Naïve Forecast

- A naïve forecast for an observation,  $Y_t$ , is the observation prior,  $Y_{t-1}$ .
- For some types of time series (e.g. **Random Walks**), a naïve forecast is the **best possible forecast one can make**.
- For almost any time series, the naïve forecast should be included as a **benchmark/baseline**. How much is your forecast better than the naïve forecast? Is it worth it?
- In the case of seasonal data, a naïve forecast could be the observation from the prior period.
  - For example, in the case of monthly data, the naïve forecast for the observation  $Y_{Jan2024}$  could be  $Y_{Jan2023}$ . In this case, we would denote the frequency,  $m=12$ , and the naïve forecast for  $Y_t$  is the observation  $m$  periods prior, or  $Y_{t-m}$ .



# Intuition and Code Behind the Naïve Forecast

```
# printing a few columns and observations from TSLA
tsla |>
  # selecting the two columns of interest
  dplyr::select(date, adjusted) |>
  # printing the last 5 observations for demo
  dplyr::slice_tail(n = 5) |>
  # printing empty space under a naive_f column
  dplyr::mutate(naive_f = '---')
```

```
## # A tibble: 5 × 3
##   date          adjusted naive_f
##   <date>          <dbl> <chr>
## 1 2022-10-01      228. ---
## 2 2022-11-01      195. ---
## 3 2022-12-01      123. ---
## 4 2023-01-01      173. ---
## 5 2023-02-01      197. ---
```

# Intuition and Code Behind the Naïve Forecast

Go to [www.menti.com/alqezoh8oh13](https://www.menti.com/alqezoh8oh13)

I know how to compute the naive forecast in R



0	0
Yes	No



# Measures of Forecast Accuracy

- The measures of accuracy we will discuss all deal with the difference between the **actual observed value** ( $Y_t$ ) and the **forecasted value** ( $F_t$ ) for time  $t$ .
- In order to measure forecast accuracy, we assume we have  $m$  actual values available, thus we have  $Y_{t+1}, Y_{t+2}, \dots, Y_{t+m}$  and forecasts  $F_{t+1}, F_{t+2}, \dots, F_{t+m}$ .
- This is important because we will be averaging the forecast errors over  $m$ .
- **Forecast Error:**  $e_{t+i} = Y_{t+i} - F_{t+i}$ .

# Measures of "Average" Forecast Performance

- A positive average error measure indicates that: for your  $m$  forecasts, on average your **actuals** ( $Y_{t+i}$ ) are **larger than** their corresponding **forecasted values** ( $F_{t+i}$ ), i.e., you are underestimating.
- A negative average error measure indicates that you are overestimating on average.
- If your average error (**percent**) measure is  $\sim 0$ ; you have an **unbiased** forecast.
  - An unbiased average measure on its own is meaningless; look at its variability.
- **Mean Error:**

$$ME = \frac{\sum_{i=1}^m e_{t+i}}{m}.$$

- **Mean Percentage Error:**

$$MPE = \frac{100}{m} \sum_{i=1}^m \frac{e_{t+i}}{Y_{t+i}}.$$

# Measures of "Variability" in Forecast Performance

- **Absolute Forecast Error:**

$$|e_{t+i}| = |Y_{t+i} - F_{t+i}|.$$

- **Squared Forecast Error:**

$$(e_{t+i})^2 = (Y_{t+i} - F_{t+i})^2.$$

- **Mean Absolute Error:**

$$MAE = \frac{\sum_{i=1}^m |e_{t+i}|}{m}.$$

- **Root Mean Squared Error:**

$$RMSE = \sqrt{\frac{\sum_{i=1}^m (e_{t+i})^2}{m}}.$$

# Measures of "Relative" Forecast Performance

- **Mean Absolute Percentage Error:**

$$MAPE = \frac{100}{m} \sum_{i=1}^m \frac{|e_{t+i}|}{m}.$$

- **Relative Mean Absolute Error:**

$$RelMAE = \frac{\sum_{i=1}^m |e_{t+i}|}{\sum_{i=1}^m |Y_{t+i} - Y_{t+i-1}|}.$$

- **Thiel's U:**

$$U = \sqrt{\frac{\sum_{i=1}^m (e_{t+i})^2}{\sum_{i=1}^m (Y_{t+i} - Y_{t+i-1})^2}}.$$

# Computing these Measures in R via the forecast

```
tsla_df =  
  tsla |>  
  # keeping relevant columns so df prints nicely  
  dplyr::select( symbol, date, year, month, adjusted) |>  
  # the naive forecast =  $Y_{t-1} = \text{lag}(Y_t, 1)$   
  dplyr::mutate( naive_f = dplyr::lag(adjusted, n =1) )  
  
# printing the first 3 rows  
tsla_df |> dplyr::slice_head(n = 3)
```

```
## # A tibble: 3 × 6  
##   symbol date      year month adjusted naive_f  
##   <chr>  <date>      <fct> <ord>      <dbl>   <dbl>  
## 1 tsla   2020-01-01 2020  Jan        43.4    NA  
## 2 tsla   2020-02-01 2020  Feb        44.5    43.4  
## 3 tsla   2020-03-01 2020  Mar        34.9    44.5
```

```
# computing the accuracy metrics via the forecast pkg  
forecast::accuracy(  
  # we start at row 2 since first fct is NA  
  object = tsla_df$naive_f[2:nrow(tsla_df)],  
  x = tsla_df$adjusted[2:nrow(tsla_df)]  
)
```

```
##  
## Test set 4.146991 40.55307 32.06263 1.861465 16.42575
```

# Computing these Measures in R with dplyr with dplyr

```
tsla_df |>
  dplyr::mutate(
    error = adjusted - naive_f, # actual - forecast
    perc_error = error / adjusted # error/actual
  ) |>
  dplyr::group_by(symbol) |> # not needed but it would allow you to compute for multiple TS
  dplyr::summarise(
    # measures of "average" forecast performance
    me = mean( error, na.rm = T ),
    mpe = 100* mean( perc_error, na.rm = T ),
    # measures of "variability" in forecast performance
    mae = mean( abs(error), na.rm = T ),
    mape = 100* mean( abs(perc_error), na.rm = T ),
    rmse = mean( error^2, na.rm = T ) |> sqrt()
  )
```

```
## # A tibble: 1 × 6
##   symbol    me    mpe    mae    mape    rmse
##   <chr>  <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 tsla    4.15  1.86  32.1  16.4  40.6
```



# Some Practical Insights

**Main Insight(s):** (Edit below)

- The **naive forecast** is .... of the series; thus, the forecast error is the ...
- In general, if the  $MAE \approx |ME|$ , then we conclude that ... . If we were using a naive forecast in such a case, then we can also conclude that ...
- Irrespective of the forecasting method, the MPE and MAPE are useful/valid if and only if ....

# The Prediction Interval

# Prediction Intervals $\neq$ Confidence Intervals

- Prediction intervals and confidence intervals are **not** the same.
- A **prediction interval** is an interval associated with a **random variable yet to be observed, with a specified probability of the random variable lying within the interval.**
  - For example, I might give an 80% interval for the forecast of GDP in 2024. The actual GDP in 2024 should lie within the interval with probability 0.8.
- A **confidence interval** is an interval associated with **a parameter** (e.g., the mean of a random variable) ... The parameter is assumed to be non-random but unknown, and the confidence interval is computed from data. Because the data are random, the interval is random ... That is, with a large number of repeated samples, 95% of the intervals would contain the true parameter if you built a 95%.
- **Prediction intervals** are wider than confidence intervals since it includes the variance of  $\epsilon$  (the error in our predictions).

# Point vs Interval Forecasts

- **Point Forecasts:** future observations for which we report a single forecast observation.
- **Interval Forecast:** a range of values that are reported to forecast an outcome.


If we assume the forecast errors follow a Normal Distribution, an approximate  $100(1 - \alpha)$  prediction interval can be computed as follows:

$$\hat{F}_t \pm Z^* * SD,$$

where:

- $\hat{F}_t$  forecast at time  $t$ .
- The RMSE can be used as an estimate of the standard deviation of the forecast errors.
- $Z^*$  is the quantile corresponding to  $100(1 - \frac{\alpha}{2})$ .

# Recall: Standard Normal Distribution



# Prediction Intervals for TSLA: "By Hand"

```
tsla_df
```

```
## # A tibble: 38 × 6
##   symbol date       year month adjusted naive_f
##   <chr> <date>     <fct> <ord>     <dbl>   <dbl>
## 1 tsla  2020-01-01  2020  Jan      43.4    NA
## 2 tsla  2020-02-01  2020  Feb      44.5    43.4
## 3 tsla  2020-03-01  2020  Mar      34.9    44.5
## 4 tsla  2020-04-01  2020  Apr      52.1    34.9
## 5 tsla  2020-05-01  2020  May      55.7    52.1
## 6 tsla  2020-06-01  2020  Jun      72.0    55.7
## 7 tsla  2020-07-01  2020  Jul      95.4    72.0
## 8 tsla  2020-08-01  2020  Aug     166.    95.4
## 9 tsla  2020-09-01  2020  Sep     143.   166.
## 10 tsla 2020-10-01  2020  Oct     129.   143.
## # ... with 28 more rows
```

# Prediction Intervals for TSLA: "By Hand"

```
# 95% prediction_interval
alpha = 0.95
curve_prob = 1 - ((1-alpha)/2)


# adding the error since it was not saved to tsla_df
tsla_df = tsla_df |> dplyr::mutate( error = adjusted - naive_f )

# recomputing the rmse and saving it to an object titled rmse
tsla_df |>
  dplyr::group_by(symbol) |>
  dplyr::summarise(rmse = error^2 |> mean(na.rm = T) |> sqrt() ) |>
  dplyr::pull(rmse) -> # pull rmse value from tibble (i.e., convert to tibble/vec)
  rmse

# computing the prediction intervals for our data
tsla_df =
  tsla_df |>
  dplyr::mutate(
    pi_l = naive_f - (abs(qnorm(curve_prob))*rmse),
    pi_u = naive_f + (abs(qnorm(curve_prob))*rmse)
  )
```

# Prediction Intervals for TSLA: "By Hand"

```
tsla_df |>
  ggplot2::ggplot(ggplot2::aes(x = date)) +
  # Plotting the actual in black
  ggplot2::geom_point(ggplot2::aes(y = adjusted)) +
  ggplot2::geom_line(ggplot2::aes(y = adjusted)) +
  # Plotting the forecast in darkgray
  ggplot2::geom_line(ggplot2::aes(y = naive_f), color =
  # Plotting the PI in red
  ggplot2::geom_ribbon(
    ggplot2::aes(ymin = pi_l, ymax = pi_u, fill = '95%
  ) +
  ggplot2::theme_bw() +
  ggplot2::theme(legend.position = 'none')
```





# Recap

# Summary of Main Points

By now, you should be able to do the following:

- Compute the nonseasonal naive forecast.
- Apply and interpret measures of forecast accuracy.
- Interpret prediction intervals for a simple forecast.

# Things to Do to Prepare for Our Next Class

- Go over your notes and read through [Chapter 2 of our reference book](#).
- **Potential Practice Problems:**
  - Extract data using either the `tq_get()` ([tidyquant package](#)) or the `covid19()` ([COVID19 package](#)), and compute the transformations using a manual (i.e., Excel) approach and R.
  - **Reference Book Example:** For the Means approaches in Example 2.7 (P.49), use R to compute the 7 error forecasting metrics (data available [here](#)).
  - **Reference Book Exercise 2.12:** Compute the forecast errors for the naive forecast.
- Complete [Assignment 05](#) on Canvas.