

ISA 444: Business Forecasting

07: Visualizing Many Time-Series

Fadel M. Megahed, PhD

Professor
Farmer School of Business
Miami University

 @FadelMegahed

 fmegahed

 fmegahed@miamioh.edu

 Automated Scheduler for Office Hours

Spring 2025

Quick Refresher of Last Class

- ✓ Describe and compute centered moving averages
- ✓ Estimate trend-cycle via moving averages
- ✓ Perform classical decomposition (trend-cycle, seasonal, residual/remainder)
- ✓ Understand STL / MSTL as alternatives to classical decomposition

Learning Objectives for Today's Class

- Explain the differences between wide vs. long format
- Use `seaborn` to plot multiple time-series
- Convert a data set to Nixtla's long format (`unique_id`, `ds`, `y`)
- Use `UtilsForecast` to visualize multiple series

Wide Vs. Long Format

Wide Format

Date	Temperature (°F)	Muggy Days	Rainy/Snowy Days
2024-01-01	31	0.0	7
2024-02-01	34	0.0	6
2024-03-01	44	0.0	9
2024-04-01	54	0.1	10
2024-05-01	64	3.4	12
2024-06-01	72	12.3	11
2024-07-01	76	18.2	11
2024-08-01	74	15.5	9
2024-09-01	67	7.0	7
2024-10-01	55	0.9	7
2024-11-01	45	0.0	7
2024-12-01	35	0.0	8

Characteristics of Wide TS Data

- Each row represents a single observation
- Each column represents a different time series
- Easy to read and understand (but not appropriate for analysis if you are using the [nixtlaverse](#) group of Python packages)

Long Format

Date	Variable	Value
2024-01-01	Muggy Days	0.0
2024-01-01	Rainy/Snowy Days	7.0
2024-01-01	Temperature (°F)	31.0
2024-02-01	Muggy Days	0.0
2024-02-01	Rainy/Snowy Days	6.0
2024-02-01	Temperature (°F)	34.0
2024-03-01	Muggy Days	0.0
2024-03-01	Rainy/Snowy Days	9.0
2024-03-01	Temperature (°F)	44.0
2024-04-01	Muggy Days	0.1
2024-04-01	Rainy/Snowy Days	10.0
2024-04-01	Temperature (°F)	54.0

Characteristics of Long TS Data

- Observations are now split into multiple rows
- Variable ids/labels are stored in a single column, and their corresponding values are stored in another column
- Easy to analyze and visualize (especially with the [nixtlaverse](#) group of Python packages)

A Visual Comparison

**Wide Format DataFrame
(Average Monthly Weather Metrics in Cincinnati)**

Date	Temperature (°F)	Muggy Days	Rainy/Snowy Days
2024-01-01	31	0.0	7
2024-02-01	34	0.0	6
2024-03-01	44	0.0	9
2024-04-01	54	0.1	10
2024-05-01	64	3.4	12
2024-06-01	72	12.3	11
2024-07-01	76	18.2	11
2024-08-01	74	15.5	9
2024-09-01	67	7.0	7
2024-10-01	55	0.9	7
2024-11-01	45	0.0	7
2024-12-01	35	0.0	8

Python pandas function:
`long_data = wide_data.reset_index().melt(id_vars=['Date'], var_name='Metric', value_name='Value')`

Class Activity: From Wide to Long Format

Description	Starter Code	Notes
-------------	--------------	-------

In this activity, you will extract the adjusted closing prices of five stocks (AAPL, MSFT, GOOGL, AMZN, TSLA) from Yahoo Finance. Once you extract the data, you should **report** the following:

- **Summary statistics** of the **adjusted closing prices for each stock**.
- **Convert the data** from wide **to long format**, and **report the shape** of the long format data.
- **Save** the long format data to a CSV file.

Hint: Read and convert the column names prior to converting the data to long format.

Visualizing Multiple Time-Series Using Seaborn

Recall: Seaborn's relplot Function

[seaborn.objects.Plot](#)

[seaborn.objects.Dot](#)

[seaborn.objects.Dots](#)

[seaborn.objects.Line](#)

[seaborn.objects.Lines](#)

[seaborn.objects.Path](#)

[seaborn.objects.Paths](#)

[seaborn.objects.Dash](#)

[seaborn.objects.Range](#)

[seaborn.objects.Bar](#)

[seaborn.objects.Bars](#)

[seaborn.objects.Area](#)

seaborn.relplot

✓ `seaborn.relplot(data=None, *, x=None, y=None, hue=None, size=None, style=None, units=None, weights=None, row=None, col=None, col_wrap=None, row_order=None, col_order=None, palette=None, hue_order=None, hue_norm=None, sizes=None, size_order=None, size_norm=None, markers=None, dashes=None, style_order=None, legend='auto', kind='scatter', height=5, aspect=1, facet_kws=None, **kwargs)`

Figure-level interface for drawing relational plots onto a FacetGrid.

This function provides access to several different axes-level functions that show the relationship between two variables with semantic mappings of subsets. The `kind` parameter selects the underlying axes-level function to use:

- `scatterplot()` (with `kind="scatter"`; the default)
- `lineplot()` (with `kind="line"`)

Extra keyword arguments are passed to the underlying function, so you should refer to the documentation for each to see kind-specific options.


Seaborn's relplot Function with Wide Data

```
import datetime as dt
import yfinance as yf
import pandas as pd
import seaborn as sns

# Download the stock data for the following companies
stock_data = (
    yf.download(
        ['AAPL', 'MSFT', 'GOOGL', 'AMZN', 'TSLA'],
        start='2020-01-01',
        end= (dt.datetime.now().date() - dt.timedelta(days=1))
    )
    # Keep only the adjusted closing price
    [['Adj Close']].reset_index()
)

# Overwrite the multi-index column names w/ single level
stock_data.columns = (
    ['Date', 'AAPL', 'MSFT', 'GOOGL', 'AMZN', 'TSLA'] )


# Plot the closing prices
sns.relplot(
    data=stock_data, kind='line', palette='Paired'
)
```



Seaborn's relplot Function with Long Data



Seaborn's relplot Function with Long Data (Facets)



Activity: Reflect on the Previous Seaborn Plots

- Whether you are plotting multiple lines in a single plot or using facets, the `relplot` function is quite versatile. However, this approach is only suitable for a few (in my opinion ≤ 9 time series). **So what options, do we have if we have more than 9 time series?**
- I think there are **three alternative charting approaches** (I am not talking about specific libraries). **Can you guess what they are?**
- In the next three minutes, edit the bullet points below to reflect the three alternative charting approaches.
 - ...
 - ...
 - ...

Advanced Visualizations with Seaborn

Approach 1: Sample

- **Sample** a subset of the time series to plot. This approach is useful when you have a large number of time series and you want to visualize a **representative sample**.
- **Key Point:** The sample should be **representative** of the entire data set.

Approach 1: Sample (Code Example)

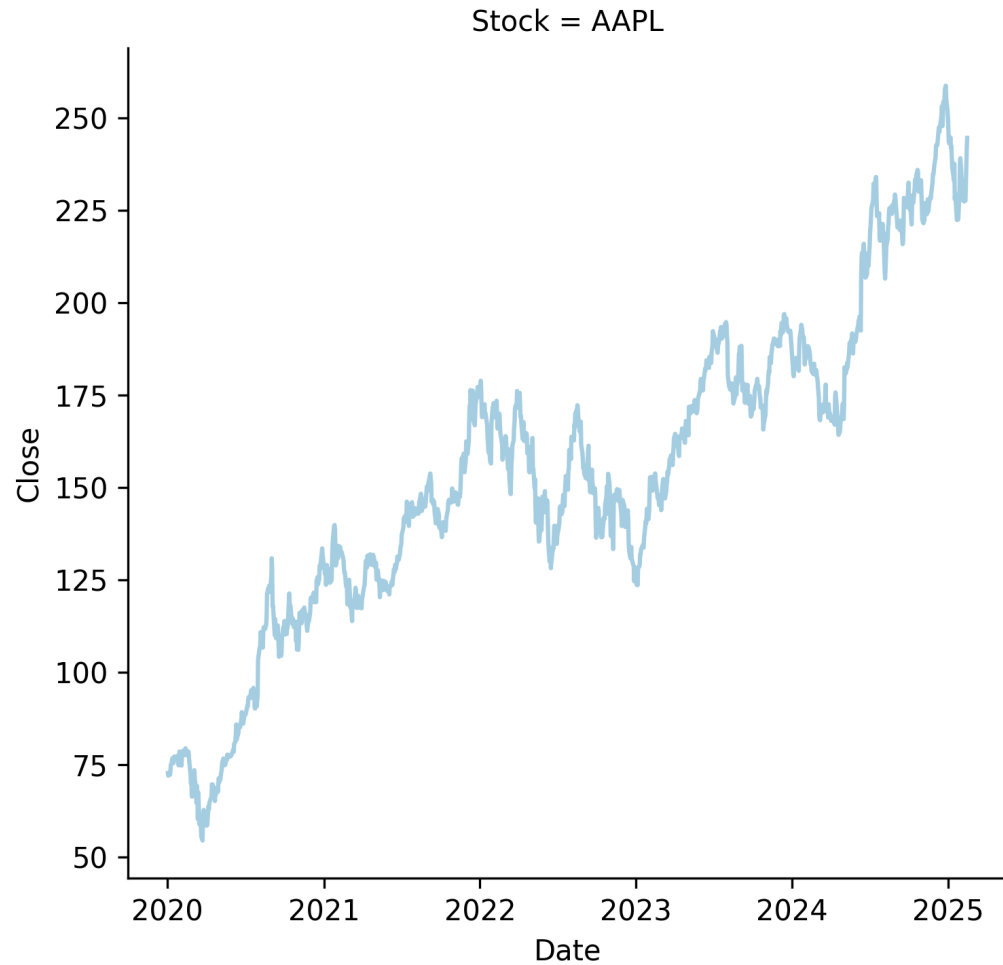
```
import pandas as pd
import seaborn as sns
import random

random.seed(2025) # for reproducibility

# Start with the long format data and sample two stocks
sampled_stocks = stock_data['Stock'].unique().tolist()
sampled_stocks = random.sample(sampled_stocks, 2)

# Plot the adjusted closing prices
fig = sns.relplot(
    data=stock_data.query('Stock in @sampled_stocks'),
    x='Date', y='Close', kind='line', legend=False, hue='Stock', col='Stock',
    palette='Paired', col_wrap=2, facet_kws={'sharey': False, 'sharex': False}
)
plt.tight_layout() # improves title visibility
```

Approach 1: Sample (Result)



Approach 2: Animated Plots

- **Animate** the time series data. This approach is useful when you have a large number of time series and you want to visualize all of them.
- **Key Point:** Animated plots can be **interactive** and **engaging**.

Approach 2: Animated Plots (Code Example)

```
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
import imageio.v2 as imageio


# Start with long format data
stocks = stock_data['Stock'].unique().tolist()

image_paths = []

for stock in stocks:
    sns.relplot(
        data=stock_data.query('Stock == @stock'), kind='line',
        x='Date', y='Close', color='black',
        height=4, aspect=3
    )
    plt.title(f"Adjusted Closing Price of {stock} (2020-2025)", fontsize=14)
    plt.xlabel("Date", fontsize=12)
    plt.ylabel("Adjusted Closing Price", fontsize=12)
    plt.tight_layout()
    plt.savefig(f'../../figures/{stock}_animated_plot.png')
    image_paths.append(f'../../figures/{stock}_animated_plot.png')

# Create a GIF from the images
images = [imageio.imread(path) for path in image_paths]
imageio.mimsave('../../figures/animated_stock_lineplot.gif', images, fps=0.25)
```

Approach 2: Animated Plots (Result)



Approach 3: Spaghetti Plot

- **Plot all time series** on a single plot. This approach is useful when you have a large number of time series and you want to visualize all of them.
- **Key Points:**
 - Use **light gray** for the lines to **de-emphasize** individual time series.
 - Use **bold colors** for the lines of **specific time series (or summary statistics across all time series)** to **emphasize** them.

Approach 3: Spaghetti Plot (Code Example)


```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# for each stock, plot the closing price over time as a light gray line
for stock, group in stock_data.groupby("Stock"):
    ax = sns.lineplot(data=group, x='Date', y='Close', color='lightgray', alpha=0.5)

# Calculate and overlay percentiles across all time series for each date
quantiles = stock_data.groupby('Date')['Close'].quantile([0.05, 0.5, 0.95]).unstack()

# Plot the median, 5th, and 95th percentiles
sns.lineplot(data=quantiles, x=quantiles.index, y=0.5, color='black', label='Median', ax = ax)
sns.lineplot(data=quantiles, x=quantiles.index, y=0.05, color='blue', label='5%', ax=ax)
sns.lineplot(data=quantiles, x=quantiles.index, y=0.95, color='red', label='95%', ax=ax)
```

Approach 3: Spaghetti Plot (Result)



Nixtla's Long Format

Class Activity: Convert Data to Nixtla's Format

- **Data:** This COVID-19 data set contains the daily cumulative number of confirmed cases for each county.
- **Objectives:**
 - Read the data set.
 - Filter the data to include only the 88 counties in Ohio, and dates from 2020-04-01 to 2022-12-31.
 - Convert the data to Nixtla's long format (`unique_id`, `ds`, `y`), where the:
 - `unique_id` column is used to identify each county.
 - `ds` column is used to represent the date.
 - `y` column is used to represent the cumulative number of confirmed cases.
 - Use the `plot_series` method from the `UtilsForecast` to visualize the data. See [here](#) to learn about how to import and [here](#) to see the arguments of the `plot_series` method.

Recap

Summary of Main Points

By now, you should be able to do the following:

- Explain the differences between wide vs. long format
- Use `seaborn` to plot multiple time-series
- Convert a data set to Nixtla's long format (`unique_id`, `ds`, `y`)
- Use `UtilsForecast` to visualize multiple series



Review and Clarification



- **Class Notes:** Take some time to revisit your class notes for key insights and concepts.
- **Zoom Recording:** The recording of today's class will be made available on Canvas approximately 3-4 hours after the session ends.
- **Questions:** Please don't hesitate to ask for clarification on any topics discussed in class. It's crucial not to let questions accumulate.