

ISA 444: Business Forecasting

08: Nonseasonal Smoothing and Forecasting

Fadel M. Megahed, PhD

Endres Associate Professor
Farmer School of Business
Miami University

 @FadelMegahed

 fmegahed

 fmegahed@miamioh.edu

 Automated Scheduler for Office Hours

Spring 2023

Recap: Lab 01 - Tidying the Data

```
macro = tidyquant::tq_get(  
  # macro indicators of interest  
  x = c('UNRATE', 'GNP', 'RHORUSQ156N'),  
  from = '2018-01-01', to = '2023-02-12',  
  # FRED Data --> set get to 'economic.data'  
  get = 'economic.data'  
)  
  
stocks_crypto = tidyquant::tq_get(  
  # you can pull stocks and crypto from  
  # the Yahoo Finance API with a single call  
  x = c("PM", "UPS", "SYK", "PCAR",  
        'BTC-USD', 'ETH-USD'),  
  from = '2018-01-01', to = '2023-02-12'  
) |> # three columns with same col names as macro  
  dplyr::select(symbol, date, adjusted) |>  
  # rename adjusted to price to match macro tibble  
  dplyr::rename(price = adjusted)  
  
# stacking/putting both tibbles on top of each other  
all_data = dplyr::bind_rows(stocks_crypto, macro)
```

```
tail(all_data, n = 12)
```

```
## # A tibble: 12 × 3  
##   symbol      date      price  
##   <chr>      <date>    <dbl>  
## 1 RHORUSQ156N 2020-01-01  65.3  
## 2 RHORUSQ156N 2020-04-01  67.9  
## 3 RHORUSQ156N 2020-07-01  67.4  
## 4 RHORUSQ156N 2020-10-01  65.8  
## 5 RHORUSQ156N 2021-01-01  65.6  
## 6 RHORUSQ156N 2021-04-01  65.4  
## 7 RHORUSQ156N 2021-07-01  65.4  
## 8 RHORUSQ156N 2021-10-01  65.5  
## 9 RHORUSQ156N 2022-01-01  65.4  
## 10 RHORUSQ156N 2022-04-01  65.8  
## 11 RHORUSQ156N 2022-07-01  66  
## 12 RHORUSQ156N 2022-10-01  65.9
```

Recap: Lab 01 - Naive Forecasting and Errors

```
# grouping by symbol to ensure calculations
# are done for each symbol separately
all_data = all_data |>
  dplyr::group_by(symbol)

all_data =
  all_data |> # now grouped
  # mutate creates new columns
  dplyr::mutate(
    # dplyr::lag since input is num vec
    naive_o = dplyr::lag(price),
    # forecasting error: actual - forecast
    error_o = price - naive_o,
    # percent error = 100*error/actual
    pe_o = 100*error_o/price
  )

all_data |>
  # summarise used --> each calc per group
  dplyr::summarise(
    # mean error per symbol
    me_o = mean(error_o, na.rm = T),
    # mean absolute percent error per symbol
    mape_o = mean(abs(pe_o), na.rm = T)
  )
```

From the printout below, **note the following**:

- Number of rows = number of symbols
- Number of columns = 1 col per each grouping variable + 1 col per each calc in summarise

```
## # A tibble: 9 × 3
##   symbol      me_o mape_o
##   <chr>      <dbl> <dbl>
## 1 BTC-USD    4.35    2.58
## 2 ETH-USD    0.397    3.49
## 3 GNP       301.     2.44
## 4 PCAR       0.0256    1.24
## 5 PM         0.0182    1.16
## 6 RHORUSQ156N 0.0895    0.676
## 7 SYK        0.0896    1.31
## 8 UNRATE     -0.01     5.61
## 9 UPS        0.0615    1.28
```

Extension: Lab 01 - Let Us Talk about Nesting

```
# the nest function creates column-lists  
# and 1 row per group  
tidyr::nest(all_data |> dplyr::select(symbol, date, price) )
```

```
## # A tibble: 9 × 2  
## # Groups:   symbol [9]  
##   symbol      data  
##   <chr>      <list>  
## 1 PM        <tibble [1,287 × 2]>  
## 2 UPS       <tibble [1,287 × 2]>  
## 3 SYK       <tibble [1,287 × 2]>  
## 4 PCAR      <tibble [1,287 × 2]>  
## 5 BTC-USD   <tibble [1,869 × 2]>  
## 6 ETH-USD   <tibble [1,869 × 2]>  
## 7 UNRATE    <tibble [61 × 2]>  
## 8 GNP       <tibble [19 × 2]>  
## 9 RHORUSQ156N <tibble [20 × 2]>
```

Extension: Lab 01 - Let Us Talk about Nesting

```
# the nest function creates column-lists
# and 1 row per group
tidyr::nest(all_data |> dplyr::select(symbol, date, price) ) |>
  # combine mutate with purrr::map
  # map transform their input by applying a function to each list
  # returning an object of the same length as the input.
  # map() always returns a list.
  dplyr::mutate(
    # create a date column-list
    date = purrr::map(.x = data, .f = magrittr::extract2, 'date'),
    # create a price column-list
    price = purrr::map(.x = data, .f = magrittr::extract2, 'price'),
    # create a column list of naive forecasts
    naive_fit = purrr::map(
      .x = price, .f = dplyr::lag
    ),
    # create a column list of accuracy metrics
    # we used map2 since we wanted to apply the function to two lists
    acc_metrics = purrr::map2(
      .x = naive_fit, .y = price,
      .f = forecast::accuracy
    ),
    # using map_dbl to have a singular value in each column
    me = purrr::map_dbl(.x = acc_metrics, .f = magrittr::extract2, c(
    mae = purrr::map_dbl(.x = acc_metrics, .f = magrittr::extract2, c(
    mape = purrr::map_dbl(.x = acc_metrics, .f = magrittr::extract2, c(
  )
```

```
## # A tibble: 9 × 9
## # Groups:   symbol [9]
##   symbol      data      date      pri
##   <chr>      <list>   <list>   <li
## 1 PM         <tibble> <date>   <db
## 2 UPS        <tibble> <date>   <db
## 3 SYK        <tibble> <date>   <db
## 4 PCAR       <tibble> <date>   <db
## 5 BTC-USD    <tibble> <date>   <db
## 6 ETH-USD    <tibble> <date>   <db
## 7 UNRATE     <tibble> <date [61]> <db
## 8 GNP        <tibble> <date [19]> <db
## 9 RHORUSQ156N <tibble> <date [20]> <db
## # ... with abbreviated variable names ^n
```

Learning Objectives for Today's Class

- Describe the benefits and drawbacks of judgmental and quantitative forecasting methods.
- Explain the difference between causal and extrapolative forecasting.
- Describe and apply smoothing with a cumulative average.
- Describe and apply forecasting with a moving average.

Extrapolative vs Causal Forecasts

Key Terms


Smoothing is usually done to help us better see patterns/trends, e.g., in time series.

- Generally smooth out the irregular roughness to see a clearer signal.
- For seasonal data, we can smooth out the seasonality so that we can identify the trend.
- Smoothing **does not provide us with a model**, but it can be a good first step in describing various components of the series.

A **filter** is sometimes used to describe a smoothing procedure. For example, **we have applied a median filter of window size 21 to smooth wearable sensors' data..**

A **forecast** is a prediction or estimate of an actual outcome expected in a future time period or for another situation

Overview of Univariate Forecasting Methods



A 10,000 foot view of forecasting techniques

Judgemental Vs Quantitative Forecasts

Judgmental Forecasting: The process of producing forecasts based on purely subjective information. The integration of subjective information may be made informally or through a structured process. The forecasts may also be obtained by aggregating the subjective forecasts of a number of individuals.

Quantitative Forecasting: Forecasting based on the application of an explicit analysis of numerical data. This kind of forecasting may be extrapolative, causal, or a blend of both.

- **Causal Forecast:** a dependent variable is forecast using explanatory variables.
- **Extrapolative:** a dependent variable is forecast using only the past values of the dependent variables. The future is "extrapolated" from the past.

Potential Drawbacks for Judgmental Forecasts

When using judgmental forecasts, forecasters may succumb to:

- An **availability bias**, when the forecaster relies too heavily on easily available and memorable information.
- The **representativeness heuristic**, when the forecaster matches a situation to a similar earlier event without taking into account its frequency of occurrence.
- The **anchoring** and **adjustment heuristic**, when the forecaster uses (anchors onto) an initial value such as the last observation and then adjusts the value to give a revised forecast.
- Over-optimism or **motivational bias** when the forecaster is motivated to bias the forecast towards a preferred state.

These biases can lead to invalid forecasts; they may lead to poor decision making, particularly when combined with overconfidence in their beliefs as to the accuracy of their forecasts.

Causal Vs Extrapolative Forecasts

Causal Forecasts

- Causal methods use data from **sources other than the series being predicted**.
- If Y is the phenomenon to forecast and X_1, X_2, \dots, X_n are the n variables we believe to be related to Y , then a causal model is one in which the forecast for Y is some function of these variables: $Y = f(X_1, X_2, \dots, X_n)$.
- Econometric models are causal models in which the relationship between Y and X_1, X_2, \dots, X_n is linear. That is:

$$Y = a_o + a_1X_1 + a_2X_2 + \dots + a_nX_n$$

for some constants a_1, a_2, \dots, a_n

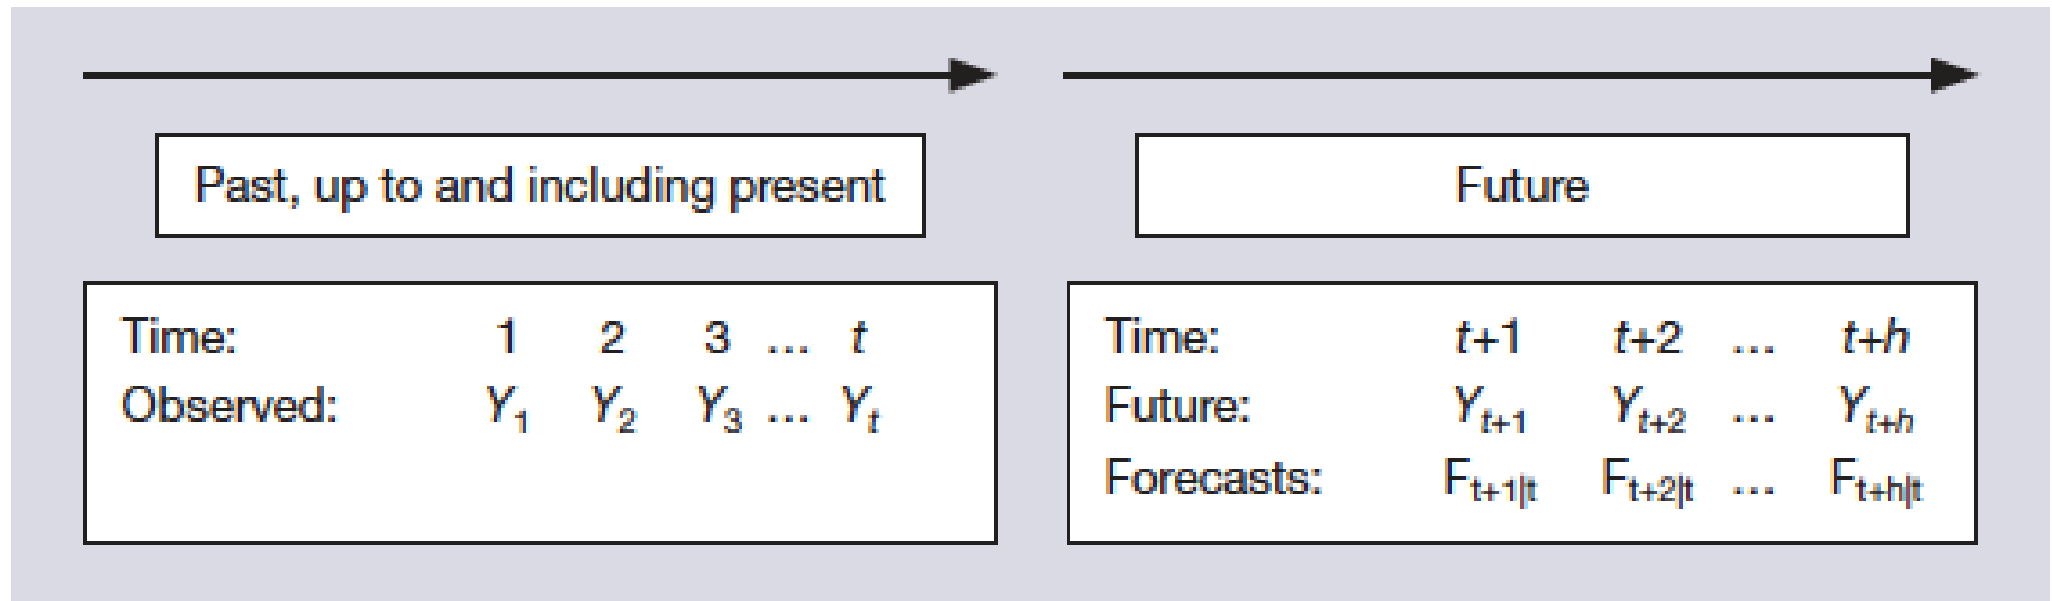
Causal Forecasts



- After watching the video, please go to [TED Ed](#).
- Answer the [five questions in the think tab](#).

Extrapolative Methods: A General Framework

Some **forecasting methods** we will discuss use the values of the series to extrapolate into the future. These **extrapolative methods** often work well for short term forecasts.



General Framework for Forecasting with a Single Series


Mean-Based Smoothing/Forecasting Techniques

Key Point: "Weak" Stationarity

Today, you are being introduced to methodologies where the time series exhibits **NO** trends and no seasonal patterns.

- *A stationary time series is one whose properties do not depend on the time at which the series is observed.*
- *Some cases can be confusing — a time series with cyclic behavior (but with no trend or seasonality) is stationary. This is because the cycles are not of a fixed length, so before we observe the series we cannot be sure where the peaks and troughs of the cycles will be.*
- *In general, a stationary time series will have no predictable patterns in the long-term. Time plots will show the series to be roughly horizontal (although some cyclic behavior is possible), with constant variance.*

Weak Stationarity: A Visual



Constancy in mean and variance

Overall Average: Almost Never Used in Practice

If you have a series that stays pretty constant over time, you could just use the overall average for smoothing. In **R**, we can use the `forecast::meanf()` to smooth using the overall mean.

```
usdc = tidyquant::tq_get(
  x = 'usdc-USD', from = '2019-01-01', to = '2023-02-01',
  periodicity = 'monthly'
)#<<

# recommendation create a ts_object of your data
u_ts = ts(data = usdc$adjusted, start = c(2019,01), frequency = 12)

u_ts
```

```
##           Jan      Feb      Mar      Apr
## 2019 1.012485 1.013673 1.001494 1.0153
## 2020 1.028374 1.004586 1.001334 1.0018
## 2021 0.999961 1.000296 0.999605 0.9998
## 2022 0.999539 0.999853 0.999447 1.0002
## 2023 1.000058 1.000057
##           Sep      Oct      Nov      Dec
## 2019 1.001386 1.004975 0.997658 1.0037
## 2020 1.000564 1.000278 1.000192 0.9998
## 2021 1.000140 0.999724 0.999456 1.0000
## 2022 1.000066 0.999996 1.000008 1.0000
## 2023
```

Overall Average: Almost Never Used in Practice

If you have a series that stays pretty constant over time, you could just use the overall average for smoothing. In **R**, we can use the `forecast::meanf()` to smooth using the overall mean.

```
# recommendation create a ts_object of your data
u_ts = ts(data = usdc$adjusted, start = c(2019,01), frequency = 12)

# using the meanf function from the forecast package
fit_overall_mean = forecast::meanf(
  y = u_ts, # time series object of interest
  h = 5, # forecast five months ahead
  level = 95 # 95% prediction interval
)

# exploring the fit_overall_mean object
class(fit_overall_mean)
names(fit_overall_mean)
```

```
## [1] "forecast"
```

```
## [1] "method" "level" "x"
## [7] "upper" "model" "lambda"
```

Some Definitions:

- *x*: The original time series
- *fitted*: Fitted values
- *residuals*: Residuals from the fitted model. That is x - fitted values
- *mean*: Point forecasts as a time series

Overall Average: Almost Never Used in Practice

If you have a series that stays pretty constant over time, you could just use the overall average for smoothing. In **R**, we can use the `forecast::meanf()` to smooth using the overall mean.

```
# recommendation create a ts_object of your data
u_ts = ts(data = usdc$adjusted, start = c(2019,01), frequency = 12)

# using the meanf function from the forecast package
fit_overall_mean = forecast::meanf( #
  y = u_ts, # time series object of interest #
  h = 5, # forecast five months ahead #
  level = 95 # 95% prediction interval #
)

# what does the model object return when printed
fit_overall_mean
```

##	Point Forecast	Lo 95	h
## Mar 2023	1.002167	0.9916601	1.01
## Apr 2023	1.002167	0.9916601	1.01
## May 2023	1.002167	0.9916601	1.01
## Jun 2023	1.002167	0.9916601	1.01
## Jul 2023	1.002167	0.9916601	1.01

Overall Average: Almost Never Used in Practice

If you have a series that stays pretty constant over time, you could just use the overall average for smoothing. In **R**, we can use the `forecast::meanf()` to smooth using the overall mean.

```
# recommendation create a ts_object of your data
u_ts = ts(data = usdc$adjusted, start = c(2019,01), frequency = 12)

# using the meanf function from the forecast package
fit_overall_mean = forecast::meanf( #
  y = u_ts, # time series object of interest #
  h = 5, # forecast five months ahead #
  level = 95 # 95% prediction interval #
)

# let us print fit_overall_mean$mean
fit_overall_mean$mean
```

```
##           Mar           Apr           May           Jun
## 2023 1.002167 1.002167 1.002167 1.002167
```

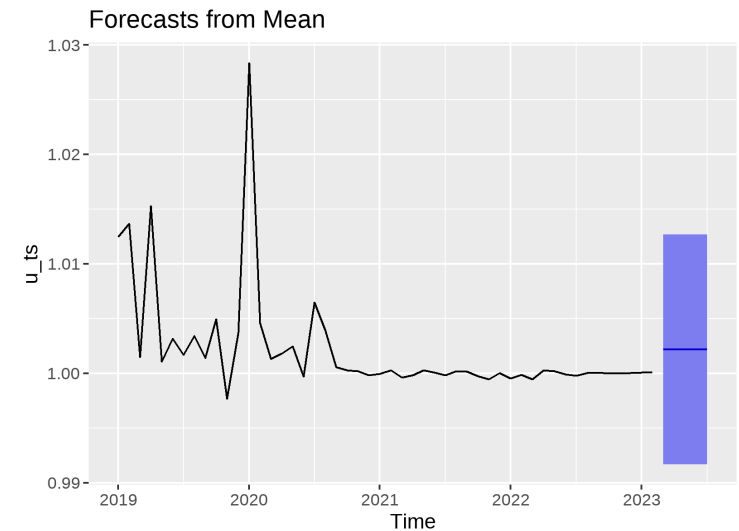
Overall Average: Almost Never Used in Practice

If you have a series that stays pretty constant over time, you could just use the overall average for smoothing. In **R**, we can use the `forecast::meanf()` to smooth using the overall mean.

```
# recommendation create a ts_object of your data
u_ts = ts(data = usdc$adjusted, start = c(2019,01), frequency = 12)

# using the meanf function from the forecast package
fit_overall_mean = forecast::meanf( #
  y = u_ts, # time series object of interest #
  h = 5, # forecast five months ahead #
  level = 95 # 95% prediction interval #
)

# let us plot what we have
forecast::autoplot(fit_overall_mean)
```



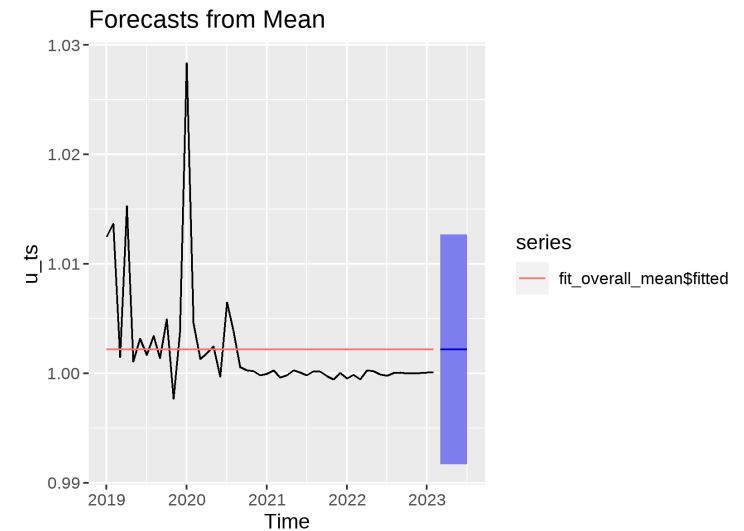
Overall Average: Almost Never Used in Practice

If you have a series that stays pretty constant over time, you could just use the overall average for smoothing. In **R**, we can use the `forecast::meanf()` to smooth using the overall mean.

```
# recommendation create a ts_object of your data
u_ts = ts(data = usdc$adjusted, start = c(2019,01), frequency = 12)

# using the meanf function from the forecast package
fit_overall_mean = forecast::meanf( #
  y = u_ts, # time series object of interest #
  h = 5, # forecast five months ahead #
  level = 95 # 95% prediction interval #
)

# let us plot what we have
forecast::autoplot(fit_overall_mean) +
  forecast::autolayer( fit_overall_mean$fitted )
```



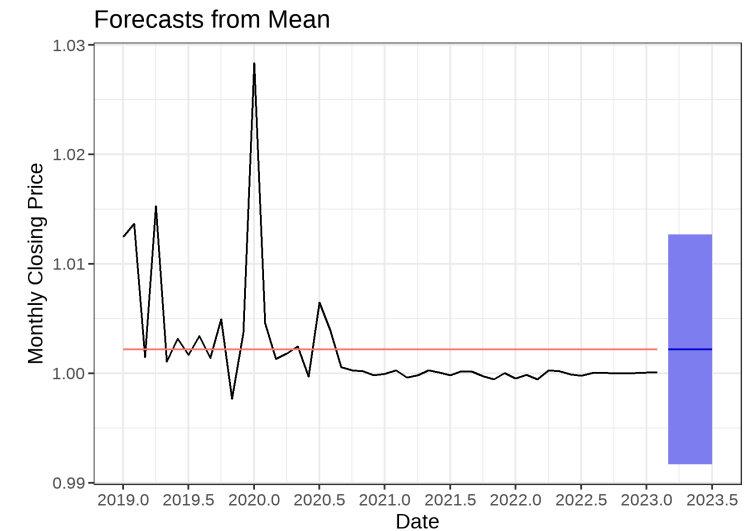
Overall Average: Almost Never Used in Practice

If you have a series that stays pretty constant over time, you could just use the overall average for smoothing. In **R**, we can use the `forecast::meanf()` to smooth using the overall mean.

```
# recommendation create a ts_object of your data
u_ts = ts(data = usdc$adjusted, start = c(2019,01), frequency = 12)

# using the meanf function from the forecast package
fit_overall_mean = forecast::meanf( #
  y = u_ts, # time series object of interest #
  h = 5, # forecast five months ahead #
  level = 95 # 95% prediction interval #
)

# let us plot what we have
forecast::autoplot(fit_overall_mean) +
  forecast::autolayer( fit_overall_mean$fitted ) +
  ggplot2::theme_bw() +
  ggplot2::theme(legend.position = 'none') +
  ggplot2::scale_x_continuous(breaks = scales::pretty_breaks(n = 10))
ggplot2::labs(x = 'Date', y = 'Monthly Closing Price')
```



Overall Average: Almost Never Used in Practice

If you have a series that stays pretty constant over time, you could just use the overall average for smoothing. In **R**, we can use the `forecast::meanf()` to smooth using the overall mean.

```
# recommendation create a ts_object of your data
u_ts = ts(data = usdc$adjusted, start = c(2019,01), frequency = 12)

# using the meanf function from the forecast package
fit_overall_mean = forecast::meanf( #
  y = u_ts, # time series object of interest #
  h = 5, # forecast five months ahead #
  level = 95 # 95% prediction interval #
)

# let us check the forecast accuracy
# We only need to input forecast object from
# (we could have also inputted the fitted and actuals)
forecast::accuracy(object = fit_overall_mean)
```

```
##                               ME          RMSE
## Training set -9.103775e-17 0.00512493
##                               ACF1
## Training set 0.2266401
```

Cumulative Average Smoothing

If you have a series that stays pretty constant over time, you could just use the **cumulative** average for smoothing. In **R**, we can use the `dplyr::cummean()` to smooth using the cumulative mean.

```
usdc = usdc |>
  dplyr::mutate(c_mean_s = dplyr::cummean(adjusted))

usdc |> dplyr::select(date, adjusted, c_mean_s)
```

```
## # A tibble: 50 × 3
##   date      adjusted c_mean_s
##   <date>      <dbl>   <dbl>
## 1 2019-01-01     1.01     1.01
## 2 2019-02-01     1.01     1.01
## 3 2019-03-01     1.00     1.01
## 4 2019-04-01     1.02     1.01
## 5 2019-05-01     1.00     1.01
## 6 2019-06-01     1.00     1.01
## 7 2019-07-01     1.00     1.01
## 8 2019-08-01     1.00     1.01
## 9 2019-09-01     1.00     1.01
## 10 2019-10-01     1.00     1.01
## # ... with 40 more rows
```

Moving Average Smoothing

If you believe recent data is more pertinent, you could just use the **moving** average for smoothing. In **R**, we can use the `zoo::rollmeanr()` to smooth using a rolling mean (moving average).

```
usdc = usdc |>
  dplyr::mutate(
    ma3 = zoo::rollmeanr(x = adjusted, k = 3, fill = NA),
    ma7 = zoo::rollmeanr(x = adjusted, k = 7, fill = NA)
  )

usdc |>
  # what I am doing in the next function can be skipped
  # printing 5 decimal points since the numbers are not varying much
  dplyr::mutate(
    dplyr::across(dplyr::where(is.numeric), ~ tibble::num(., digits = 5))
  )
  dplyr::select(
    date, adjusted, c_mean_s, ma3, ma7
  )
```

```
## # A tibble: 50 × 5
##   date      adjusted c_mean_s
##   <date>      <num:.5!> <num:.5!> <num:.5!>
## 1 2019-01-01  1.01249  1.01249  NA
## 2 2019-02-01  1.01367  1.01308  NA
## 3 2019-03-01  1.00149  1.00922  1.
## 4 2019-04-01  1.01530  1.01074  1.
## 5 2019-05-01  1.00105  1.00880  1.
## 6 2019-06-01  1.00319  1.00787  1.
## 7 2019-07-01  1.00169  1.00698  1.
## 8 2019-08-01  1.00344  1.00654  1.
## 9 2019-09-01  1.00139  1.00597  1.
## 10 2019-10-01  1.00498  1.00587  1.
## # ... with 40 more rows
```

From Smoothing to Forecasting

So Can We Use the Previous Results for Forecasting?

Comments on Previously Printed Results: (Insert below)

- **c_mean_s:**
- **ma3:**
- **ma7:**
- **Equivalence of c_mean_s with ma:**

A Live Demo to Generate Forecasts

Recap

Summary of Main Points

By now, you should be able to do the following:

- Describe the benefits and drawbacks of judgmental and quantitative forecasting methods.
- Explain the difference between causal and extrapolative forecasting.
- Describe and apply smoothing with a cumulative average.
- Describe and apply forecasting with a moving average.

Things to Do to Prepare for Our Next Class

- **Recommended:** Thoroughly read [Chapter 3.1-3.2 of our reference book](#).
- Go through the slides, examples and make sure you have a good understanding of what we have covered.
- If you are interested in additional practice problems, please consider the following problems from your textbook. To access these datasets, please click [here](#).
 - **Practice Problem:** Exercise 3.1 (RMSE is 2.54 and 2.28 for the MA 3 and MA7, respectively).
- **Required:** Complete [assignment06](#), and [assignment07](#).