

ISA 444: Business Forecasting

25: Time Series Regression

Fadel M. Megahed, PhD

Endres Associate Professor
Farmer School of Business
Miami University

 @FadelMegahed

 fmegahed

 fmegahed@miamioh.edu


 Automated Scheduler for Office Hours

Spring 2023

Quick Refresher from Last Class

- ✓ Explain the simple and multiple linear regression models and interpret the parameters.
- ✓ Interpret the sample linear regression coefficients in the language of the problem.
- ✓ Use a simple linear regression model for trend adjustment (time-series data).

Overview of Univariate Forecasting Methods



A 10,000 foot view of forecasting techniques

Learning Objectives for Today's Class

- Use a simple linear regression model for trend adjustment (time-series data).
- Interpret regression diagnostic plots.
- Create prediction intervals for individual values of the response variable.

Using Simple Linear Regression Model for Trend Adjustment (with Time Series Data)

Continuing our Example from Last Class

```
jj = astsa::jj

# Step 1: Plot the ts data (saved to object to not
p = forecast::autoplot(jj) +
  ggplot2::theme_bw() +
  ggplot2::geom_point(size = 1)

# stabilizing the variance
log_jj = log(jj)

# Step 1b: Our updated plot
p2 = forecast::autoplot(log_jj) +
  ggplot2::theme_bw() +
  ggplot2::geom_point(size = 1)

# Step 2: Extract time
year = time(log_jj)

# Step 3: Fit the regression model
reg_model = lm(log_jj ~ year)

summary(reg_model) # prints top right

anova(reg_model) # prints bottom right
```

```
##
## Call:
## lm(formula = log_jj ~ year)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.38309 -0.08569  0.00297  0.09984  0.38016
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -3.275e+02  5.623e+00  -58.25  <2e-16 ***
## year         1.668e-01  2.854e-03   58.45  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
##
## Residual standard error: 0.1585 on 82 degrees of freedom
## Multiple R-squared:  0.9766,    Adjusted R-squared:  0.9
## F-statistic: 3416 on 1 and 82 DF,  p-value: < 2.2e-16
```

```
## Analysis of Variance Table
##
## Response: log_jj
##              Df Sum Sq Mean Sq F value    Pr(>F)
## year           1 85.872  85.872  3416.5 < 2.2e-16 ***
## Residuals     82  2.061   0.025
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1
```

Continuing our Example from Last Class

What would be our forecast values log EPS for 1981?

Manual Calculations:

Recall that our regression equation from `summary(reg_model)` was:

$$\log_jj = -327.5 + 0.1668 \times \text{year}$$

- For Q1, my predicted log EPS = ...
- For Q2, my predicted log EPS = ...
- For Q3, my predicted log EPS = ...
- For Q4, my predicted log EPS = ...

R Functions:

```
# Approach (a):  
pred1981a = predict(  
  object = reg_model,  
  newdata =  
    data.frame(year = c(1981, 1981.25, 1981.5, 1981  
)  
  
# Approach (b): I prefer this approach  
pred1981b = forecast::forecast(  
  object = reg_model,  
  newdata =  
    data.frame(year = c(1981, 1981.25, 1981.5, 1981  
)
```

Why is approach (b) slightly better? (ignoring the need to load an extra package)

Approach (b) is better because:

forecast::tslm() vs. lm()

Let us examine and contrast the output from the `forecast::tslm()`, compared to what we obtained in Slide 6. There are **3** reasons for which I tend to prefer the `forecast::tslm()` (if you are willing to ignore the need to load an extra package):

Please fill after we go through the example:

- ...
- ...
- ...

Our Example with `forecast::tslm()`

```
model_tslm =  
  forecast::tslm(log_jj ~ trend)  
  
summary(model_tslm)
```

From the output, what are the intercept and slope for year? Interpret their values.

Slope and intercept from `forecast::tslm()`:

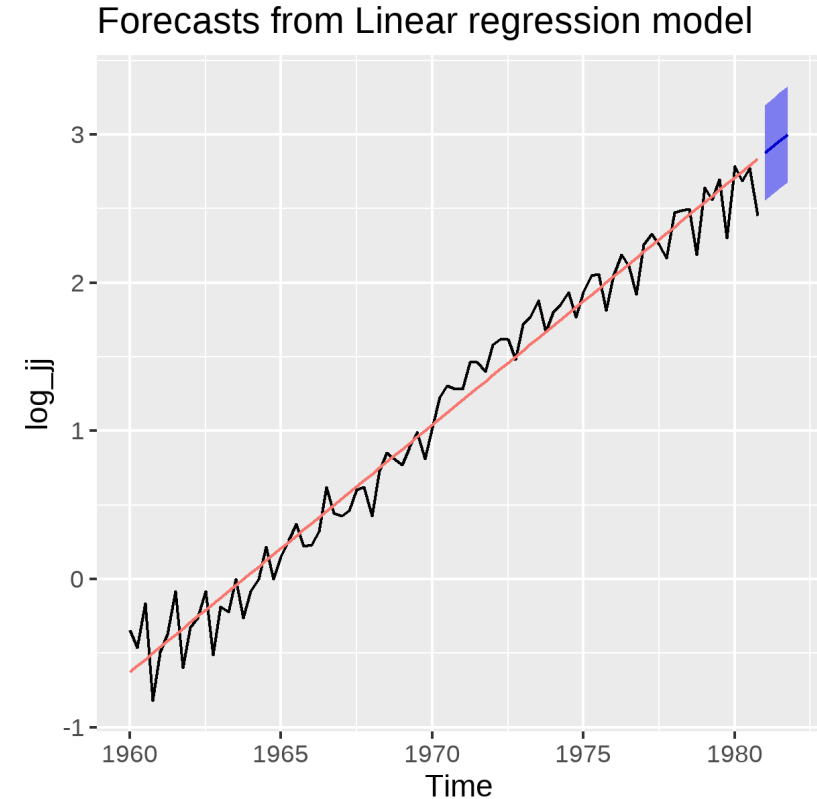
- Intercept: ...
- Interpretation: ...
- Slope: ...
- Interpretation: ...
- Difference from `lm()` output: ...

```
##  
## Call:  
## forecast::tslm(formula = log_jj ~ trend)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -0.38309 -0.08569  0.00297  0.09984  0.38016   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept) -0.6677756  0.0349073  -19.13  <2e-16   
## trend        0.0416992  0.0007134   58.45  <2e-16   
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 0.1585 on 82 degrees of freedom  
## Multiple R-squared:  0.9766,    Adjusted R-squared:  0.9754   
## F-statistic: 3416 on 1 and 82 DF,  p-value: < 2.2e-16
```

Our Example with `forecast::tslm()`

```
model_tslm =  
  forecast::tslm(log_jj ~ trend)  
  
pred_1981_tslm =  
  forecast::forecast(  
    model_tslm, h = 4, level = 95  
  )  
  
# Printing the point estimates  
pred_1981_tslm$mean  
  
# TS Plot w/ forecast, fitted values, &  
# 95% PI ( defined in forecast() )  
forecast::autoplot(pred_1981_tslm) +  
  forecast::autolayer(  
    fitted(pred_1981_tslm)  
  ) +  
  ggplot2::theme(  
    legend.position = 'none'  
  )
```

```
##           Qtr1      Qtr2      Qtr3      Qtr4  
## 1981 2.876655 2.918354 2.960053 3.001752
```



Regression Diagnostic Plots

Regression Diagnostic Plots

```
resplot <- function(res, fit, ncol = 2, nrow = 3, freq = NULL){
  ggplot2::theme_set(ggplot2::theme_bw()) # setting the theme for all ggplots to be black and white

  df = data.frame(res = res, fit = fit) # creating a df of residuals and fitted values

  # Plot 1: QQ Plot for the residuals
  qq = ggplot2::ggplot(df, ggplot2::aes(sample = res)) +
    ggplot2::stat_qq() + ggplot2::stat_qq_line() + ggplot2::labs(x = "Theoretical Values", y = "Sample")

  # Plot 2: Scatter Plot of Residuals Vs. Fitted Values
  scatter = ggplot2::ggplot(df, ggplot2::aes(x = fit, y = res)) +
    ggplot2::geom_point() + ggplot2::geom_hline(yintercept = 0) +
    ggplot2::labs(x = "Fitted Values", y = "Residuals")

  # Plot 3: Histogram of Residuals with Density Plot
  histogram = ggplot2::ggplot(df, aes(x = res)) +
    ggplot2::geom_histogram(bins = 15) + ggplot2::geom_density(alpha = 0.2, fill = "red") +
    ggplot2::geom_vline(ggplot2::aes(xintercept= mean(res)), color="red", linetype="dashed", size=1) + ggplot2::labs(x = "Residuals", y = "Count")

  # Plot 4: Time Order Plot of Residuals
  if(is.numeric(freq)){
    colorFactor = as.factor( as.numeric(row.names(df)) %% freq ) # coloring based on the modulus operation
    timeOrder = ggplot2::ggplot(df, ggplot2::aes(x = as.numeric(row.names(df)), y = res) ) +
      ggplot2::geom_line() + ggplot2::geom_point(aes(color = colorFactor)) + ggplot2::geom_hline(yintercept = 0) +
      ggplot2::labs(x = "Observation Order", y = "Residuals") + ggplot2::theme(legend.position = "none")
  }else{
    timeOrder = ggplot2::ggplot(df, ggplot2::aes(x = as.numeric(row.names(df)), y = res)) +
      ggplot2::geom_line() + ggplot2::geom_point() + ggplot2::geom_hline(yintercept = 0) +
      ggplot2::labs(x = "Observation Order", y = "Residuals")
  }

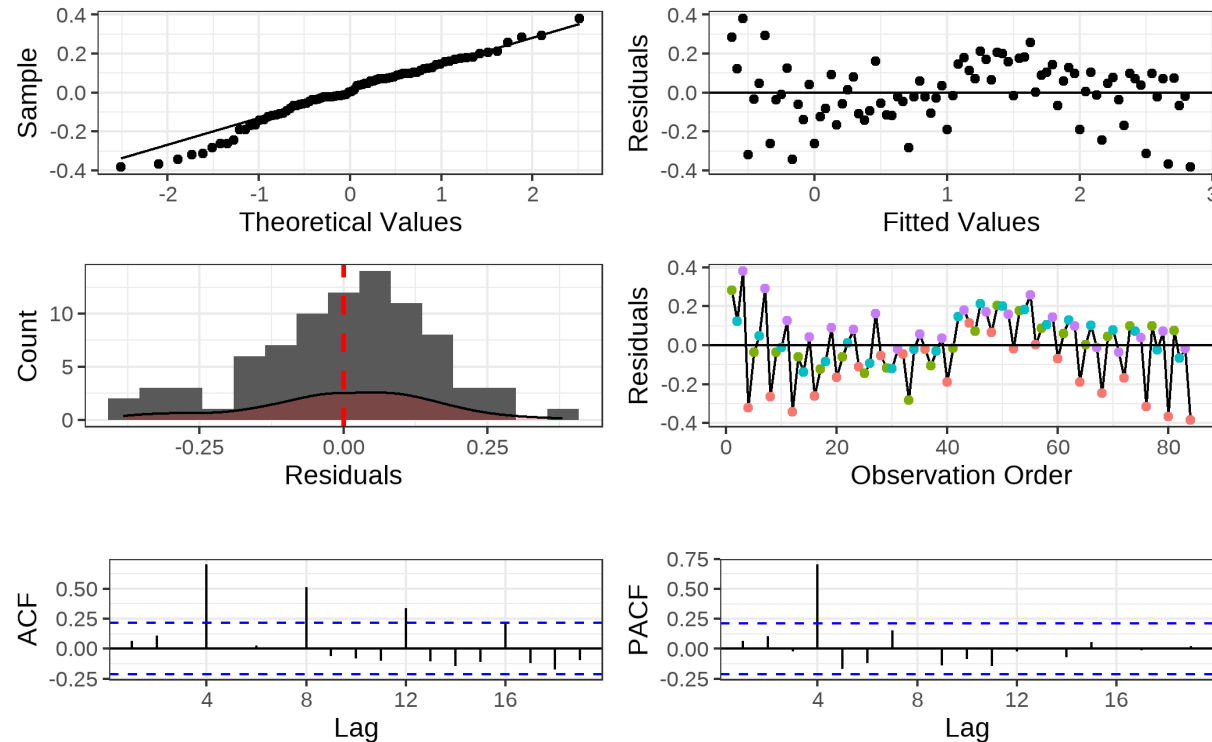
  # Plot 5: ACF Plot of Residuals
  acfPlot = forecast::ggAcf(res) + ggplot2::labs(title = "")

  # Plot 6: ACF Plot of Residuals
  pacfPlot = forecast::ggPacf(res) + ggplot2::labs(title = "")

  # Putting all 6 figures together
  ggpubr::ggarrange(qq, scatter, histogram, timeOrder, acfPlot, pacfPlot, ncol = ncol, nrow = nrow)
}
```

Applying the `resPlot()`

```
residuals_tslm = model_tslm$residuals  
fitted_tslm = model_tslm$fitted.values  
  
resplot(res = residuals_tslm, fit = fitted_tslm, freq = 4)
```



Insights from the `resPlot()`

- Assumption behind using the forecast function is that we have a model that fits well to our data.
- The diagnostics on the regression model confirmed that we do not have an excellent model since our residuals are not iid. Specifically, the:
 - The residuals **vary in magnitude as a function of the fitted values**.
 - The residuals show a **seasonal pattern**, where Q2 residuals are typically large and Q4 residuals are typically small.
 - The ACF and PACF show a **seasonal pattern in the residuals** – likely need to fit an AR model.

Recap

Summary of Main Points

By now, you should be able to do the following:

- Use a simple linear regression model for trend adjustment (time-series data).
- Interpret regression diagnostic plots.
- Create prediction intervals for individual values of the response variable.

Things to Do to Prepare for Next Class

- Go through the slides, examples and make sure you have a good understanding of what we have covered.
- Read Chapter 7 in our reference book [Principles of Business Forecasting](#).