

ISA 444: Business Forecasting

28: A 20-minute Introduction to ML for TS Data

Fadel M. Megahed, PhD

Endres Associate Professor
Farmer School of Business
Miami University

 @FadelMegahed

 fmegahed

 fmegahed@miamioh.edu

 Automated Scheduler for Office Hours

Spring 2023

Quick Refresher from Last Class

- ✓ Combine regression with ARIMA models to model a time series with autocorrelated errors.
- ✓ Use the `xreg` argument to combine ARIMA models with regression predictors.

Class Activity Solution

```
uschange = fpp2::uschange

# Solutions for the Questions
# (1a) Extrapolative forecasting using auto.arima (i.e., only the time-series for Consumption)
model2 = forecast::auto.arima(uschange[, 'Consumption'])
summary(model2)
resplot(res = model2$residuals, fit = model2$fitted, freq = 4)

# (1b) Reg with Income and Savings and ARIMA structure imposed on the error term
model3 = forecast::auto.arima(uschange[, 'Consumption'], xreg = uschange[, c('Income', 'Savings') ])
summary(model3)
resplot(res = model3$residuals, fit = model3$fitted, freq = 4)

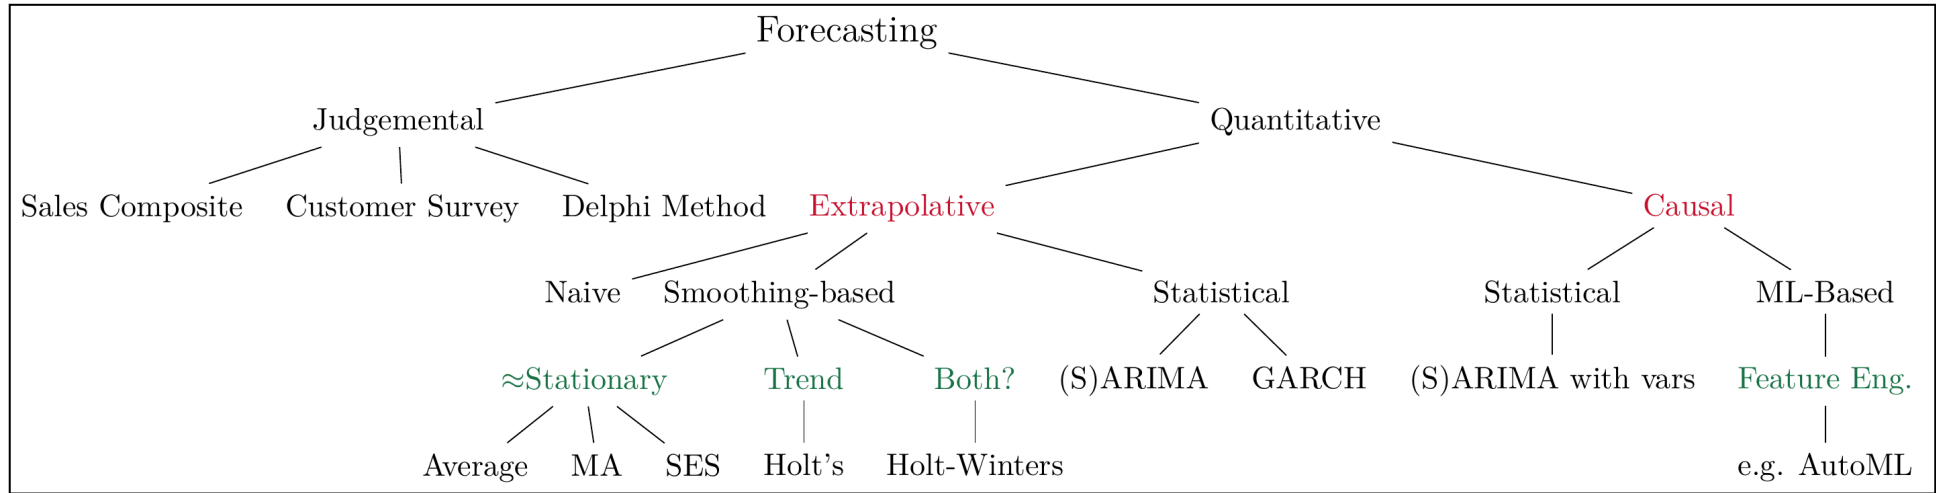
# (1c) Lm using both income and savings
model4 = lm(uschange[, 'Consumption'] ~ uschange[, 'Income'] + uschange[, 'Savings'])
summary(model4)

# (1d) Income only
model5 = lm(uschange[, 'Consumption'] ~ uschange[, 'Income'] )
summary(model5)
resplot(model5$residuals, model5$fitted.values, freq = 4)

# How to make predictions about future values (made up values)
predict(model1, newxreg = c(0.1, -0.2)) # for model1 (one predictor)
# Alternatively
forecast::forecast(model1, xreg = c(0.1, -0.2)) # for model1 (one predictor)

# For forecast models with multiple predictors, you will have to add the information as a data.matrix
forecast::forecast(model3,
  xreg = data.frame( Income = c(0.1, -0.2), Savings = c(0.1, 0.2) ) |> data.matrix() )
```

Overview of Univariate Forecasting Methods



A 10,000 foot view of forecasting techniques

Learning Objectives for Today's Class

- Explain how ML, and other advanced models, can be applied to TS data (given that we will be introducing this for **20 minutes** prior to answering questions pertaining to your final exam, this will be a very quick demo).

ML for TS Data

(An Example from Fadel's Research)

COVID Deaths in Saint Louis City, MO

```
# creating a temp file for downloading the data
temp = tempfile()

# download the file to temporary location
download.file("https://storage.covid19datahub.io/country/USA.csv.zip", temp)

# unzip and read the file
covid_tbl = unz(temp, "USA.csv") |>
  # reading the data from the CSV
  readr::read_csv() |>
  # filtering to Missouri and Illinois
  dplyr::filter(administrative_area_level_2 %in% c('Illinois', 'Missouri'))

st_louis_tbl = covid_tbl |>
  dplyr::filter(
    (administrative_area_level_2 == 'Illinois' &
     administrative_area_level_3 %in% c('Bond', 'Calhoun', 'Clinton', 'Jersey', 'Macoupin', 'Madison', 'Monroe')) |
    (administrative_area_level_2 == 'Missouri' &
     administrative_area_level_3 %in% c('Crawford', 'Franklin', 'Jefferson', 'Lincoln', 'St. Charles', 'St. Clair', 'St. Louis', 'St. Louis City', 'Wa
  )


# Aggregating the counts by day (so that we have an approximation of total numbers for st. louis)
st_louis_agg_tbl =
  st_louis_tbl |>
  # grouping by date so we can create an aggregated summation across all counties
  dplyr::group_by(date) |>
  dplyr::select(date,
    # variables to be summed across all counties in St. Louis
    confirmed, deaths, recovered, hosp, icu, vent,
    population,
    # other variables that can be included in the analysis later
    stringency_index) |>
  dplyr::summarise_at(
    dplyr::vars(confirmed, deaths, recovered, hosp, icu, vent, population),
    .funs = sum, na.rm = T
  ) |> dplyr::ungroup()

unlink(temp) # remove temp file
```

Visualizing the TS Data

```
st_louis_agg_tbl |>
  tidyr::pivot_longer(cols = c(confirmed, deaths, icu),
    names_to = 'statistic') |>
  ggplot2::ggplot(
    ggplot2::aes(x = date, y = value)
  ) +
  ggplot2::geom_line() +
  ggplot2::facet_wrap(~ statistic, scales = 'free_y', ncol = 1) +
  ggplot2::theme_bw() +
  ggplot2::scale_x_date(breaks = scales::pretty_breaks(n = 6)) +
  ggplot2::scale_y_continuous(labels = scales::comma) +
  ggplot2::labs(title = 'Plots of the Time-Series of Potential Variables for St. Louis MSA',
    caption = 'Based on data aggregated from the COVID19 DataHub',
    x = 'Date',
    y = NULL)
```


Visualizing the TS Data



Based on data aggregated from the COVID19 DataHub


Updated Time Series

```
st_louis_agg_tbl =  
  st_louis_agg_tbl |>  
  # removing anomalies  
  dplyr::filter(  
    date >= lubridate::ymd('2021-01-01') &  
    date <= lubridate::ymd('2023-03-23')) |>  
  # creating potential predictors  
  dplyr::mutate(  
    # based on https://www.nytimes.com/interactive/2021/health/coronavirus-variant-tracker.html  
    dominant_variant =  
      dplyr::case_when(  
        date < lubridate::ymd('2021-03-01') ~ 'Epsilon',  
        date < lubridate::ymd('2021-06-15') ~ 'Alpha',  
        date < lubridate::ymd('2021-12-15') ~ 'Delta',  
        date >= lubridate::ymd('2021-12-15') ~ 'Omicron'  
      ) |> forcats::as_factor(),  
    # creating a list of special holidays  
    holidays =  
      dplyr::if_else(date %in% tidyquant::HOLIDAY_SEQUENCE(start_date = min(date),  
                                                             end_date = max(date),  
                                                             calendar = 'NYSE'),  
                     true = 'yes',  
                     false = 'no') |> forcats::as_factor()  
  ) |>  
  tidyr::drop_na()
```

Visualizing the Updated TS

```
st_louis_agg_tbl |>
  tidyr::pivot_longer(cols = c(confirmed, deaths),
    names_to = 'statistic') |>
  ggplot2::ggplot(
    ggplot2::aes(x = date, y = value, color = dominant_variant)
  ) +
  ggplot2::geom_line() +
  ggplot2::facet_wrap(~ statistic, scales = 'free_y', ncol = 1) +
  ggplot2::theme_bw() +
  ggplot2::scale_x_date(breaks = scales::pretty_breaks(n = 6)) +
  scale_y_continuous(labels = scales::comma) +
  ggplot2::scale_color_brewer(palette = 'Paired') +
  ggplot2::labs(x = 'Date',
    y = 'New Daily Counts',
    title = 'Plots of Confirmed Cases and Deaths for St. Louis MSA')
```

Visualizing the Updated TS



Creating time splits for Training and Validation

```
splits = rsample::initial_time_split(st_louis_agg_tbl, prop = 0.9)

print(splits)

paste('The starting and ending dates for training are',
      splits$data[splits$in_id, 'date'] |> head(n=1) |> dplyr::pull(),
      'and',
      splits$data[splits$in_id, 'date'] |> tail(n=1) |> dplyr::pull(),
      'respectively. For the holdout data, the starting and training dates are',
      splits$data[splits$out_id, 'date'] |> head(n=1) |> dplyr::pull(),
      'and',
      splits$data[splits$out_id, 'date'] |> tail(n=1) |> dplyr::pull())
```

```
## <Training/Testing/Total>
## <730/82/812>
## [1] "The starting and ending dates for training are 2021-01-01 and 2022-12-31 respectively"
```

Training Different Time-Series Models

I have quickly trained the following three models:

- A univariate Auto ARIMA model with no xreg
- An Auto ARIMA model with confirmed, holidays (NYSE holidays) and dominant variant as our xreg
- The Prophet Model, originally developed by [Facebook](#). See the [Forecasting at Scale Paper](#) for more details.

auto.arima() with no xreg

```
library(modeltime)  
# a univariate ARIMA model using "Auto Arima" using arima_reg()  
# using the modeltime pkg this will automatically pick the weekly seasonality  
model_fit_arima =  
  modeltime::arima_reg() |>  
  parsnip::set_engine(engine = "auto_arima") |> # this requires library(modeltime)  
  parsnip::fit(deaths ~ date, data = rsample::training(splits) )
```

auto.arima() with xreg

```
# ARIMA with xreg
model_fit_arima_xreg =
  modeltime::arima_reg() |>
  parsnip::set_engine(engine = "auto_arima") |>
  # confirmed, holidays and dominant variant as our xreg
  parsnip::fit(
    deaths ~ date + confirmed + holidays + dominant_variant,
    data = rsample::training(splits)
  )
```


The Prophet Model

```
library(prophet)

model_fit_prophet =
  modeltime::prophet_reg() |>
  parsnip::set_engine(engine = "prophet") |>
  parsnip::fit(deaths ~ date + confirmed + holidays + dominant_variant,
    data = rsample::training(splits) )
```

Model Table

```
models_tbl =  
  modeltime::modeltime_table(  
    model_fit_arima,  
    model_fit_arima_xreg,  
    model_fit_prophet  
  )
```

```
models_tbl
```

```
## # Modeltime Table  
## # A tibble: 3 × 3  
##   .model_id .model      .model_desc  
##         <int> <list>    <chr>  
## 1           1 <fit[+]> ARIMA(2,2,2)(2,0,0)[7]  
## 2           2 <fit[+]> REGRESSION WITH ARIMA(0,0,4)(0,0,2)[7] ERRORS  
## 3           3 <fit[+]> PROPHET W/ REGRESSORS
```


Training Performance: Residuals

```
models_tbl |>
  modeltime::modeltime_calibrate(new_data = rsample::training(splits)) |>
  modeltime::modeltime_residuals() |>
  modeltime::plot_modeltime_residuals(.interactive = FALSE,
                                     .type = 'timeplot') +
  ggplot2::scale_x_date(breaks = scales::pretty_breaks(12)) +
  ggplot2::scale_color_brewer(palette = 'Dark2') +
  ggplot2::facet_wrap(~ .model_desc, ncol = 1, scales = 'free') +
  ggplot2::theme_bw() +
  ggplot2::theme(legend.position = 'none') +
  ggplot2::labs(
    title = 'Residuals plot for the three models based on our training data',
    subtitle = 'The residuals are large on the same day irrespective of model')
```

Training Performance: Residuals

Residuals plot for the three models based on our training data

The residuals are large on the same day irrespective of model



Statistical Tests for Residuals

```
models_tbl |>
  modeltime::modeltime_calibrate(new_data = rsample::training(splits)) |>
  modeltime::modeltime_residuals() |>
  modeltime::modeltime_residuals_test()
```

```
## # A tibble: 3 × 6
##   .model_id .model_desc      shapiro_wilk box_pierce  ljung_box  durbin_watson
##   <int> <chr>          <dbl>      <dbl>      <dbl>      <dbl>
## 1       1 ARIMA(2,2,2)(2,0,0)... 3.28e-35    0.846    0.846    2.01
## 2       2 REGRESSION WITH ARI... 6.57e-25    0.956    0.956    1.93
## 3       3 PROPHET W/ REGRESSO... 4.54e-16    0        0        0.456
```

Training Performance

```
models_tbl |>
  modeltime::modeltime_calibrate(new_data = rsample::training(splits)) |>
  modeltime::modeltime_accuracy() |>
  modeltime::table_modeltime_accuracy()
```


↕	.model_id	.model_desc	↕	.type	↕	mae	↕	mape	↕	mase	↕	smape	↕	rmse	↕	rsq
	1	ARIMA(2,2,2)(2,0,0)[7]		Fitted		6.28		0.11		0.81		0.11		11.66		1
	2	REGRESSION WITH ARIMA(0,0,4)(0,0,2)[7] ERRORS		Fitted		16.14		0.3		2.08		0.3		24.73		1
	3	PROPHET W/ REGRESSORS		Fitted		11.44		0.19		1.48		0.19		16.27		1

Training Performance: Residuals

```
models_tbl |>
  modeltime::modeltime_calibrate(new_data = rsample::testing(splits)) |>
  modeltime::modeltime_residuals() |>
  modeltime::plot_modeltime_residuals(.interactive = FALSE,
                                     .type = 'timeplot') +
  ggplot2::scale_x_date(breaks = scales::pretty_breaks(12)) +
  ggplot2::scale_color_brewer(palette = 'Dark2') +
  ggplot2::facet_wrap(~ .model_desc, ncol = 1, scales = 'free') +
  ggplot2::theme_bw() +
  ggplot2::theme(legend.position = 'none') +
  ggplot2::labs(
    title = 'Residuals plot for the three models based on our testing data')
```

Testing Performance: Residuals

Residuals plot for the three models based on our testing data



Testing Performance

```
models_tbl |>
  modeltime::modeltime_calibrate(new_data = rsample::testing(splits)) |>
  modeltime::modeltime_accuracy() |>
  modeltime::table_modeltime_accuracy()
```

↕	.model_id	.model_desc	↕	.type	↕	mae	↕	mape	↕	mase	↕	smape	↕	rmse	↕	rsq
	1	ARIMA(2,2,2)(2,0,0)[7]		Test		9.82		0.11		2.89		0.11		11.66		0.98
	2	REGRESSION WITH ARIMA(0,0,4)(0,0,2)[7] ERRORS		Test		112.53		1.29		33.15		1.28		115.24		0.91
	3	PROPHET W/ REGRESSORS		Test		17.77		0.2		5.23		0.2		19.91		0.99

Recap

Summary of Main Points

By now, you should be able to do the following:

- Explain how ML, and other advanced models, can be applied to TS data (given that we will be introducing this for **20 minutes** prior to answering questions pertaining to your final exam, this will be a very quick demo).

Things to Do to Prepare for the Final Exam

- Go through the slides, examples and make sure you have a good understanding of what we have covered.
- **Exam Setup:**
 - Q1 and Q2 interpretation of regression coefficients
 - Q3 interpretation of a residuals plot (`resPlot()`)
 - Q4 and Q5 interpretation of `tslm()` outputs
 - Q6 Interpretation of a `lm()` or `tslm()` model summary
 - Q7-Q8 interpretation of ARIMA with xreg
 - Q9-Q22 interpretations of which models to fit, autocorrelation, etc based on a plot of a time-series and its ACF
 - Q23-Q32 conceptual multiple choice questions